

Short-lead forecasting of GRACE terrestrial water storage is a noise-filtering problem: an observation-noise-aware benchmark, its consequences, and where exogenous forcing takes over

Ishaan Kejriwal¹

¹Independent researcher

Correspondence: Ishaan Kejriwal (ishaankejriwal1@gmail.com)

Abstract. Forecasts of GRACE/GRACE-FO terrestrial water storage anomalies (TWSA) are routinely benchmarked against persistence. We show that for basin-average, deseasonalized TWSA this benchmark leaves the persistence class’s own headroom unexploited: the observations are noisy, and persistence propagates the anchor month’s noise into the forecast. A per-basin AR(1)-plus-observation-noise state-space model – a three-parameter Kalman filter, not a data-assimilation system – filters the series before extrapolating. On 234 global basins over five expanding-window test folds (issue months June 2019 onward, 83 at lead 1 shrinking to 78 at lead 6 as the shifted target runs off the May 2026 record end), this filter beats damped persistence by 5.0 % and 8.8 % at leads 1–2 and per-basin ridge regression at five of six leads (Diebold–Mariano $p \leq 0.013$; lead 4 is not significant, $p = 0.078$). An ablation identifies the mechanism: with the observation-noise variance forced to zero the same model class falls behind damped persistence at leads 2–6, so the margin is the noise-filtering term itself, not the autoregressive structure. Much of the published skill of learned models over persistence is therefore attainable by three-parameter state-space filtering alone. Exogenous forcing is what beats this benchmark: ERA5-Land predictors add 4.6 % at lead 1 and nothing detectable beyond it, while extending the forcing window to twelve months adds a further 3.3–5.0 % at leads 1–3 – read *linearly*: a ridge over the flattened history beats every sequence model we trained on it. A head-to-head comparison with the GRACE-FCast system of Li and Kusche (2026) on 227 matched basins locates the classic initial-conditions-versus-forcing crossover on satellite-observed TWSA: our filtering-based system wins lead 1 by 20 %, the systems are not detectably different at lead 2, and their forcing-driven system wins leads 3–6 by 13–30 %. Finally, we characterize the cross-basin signal in observed GRACE TWSA under null-model controls. At the linear tier it is a null result: one correlation-selected neighbor adds +0.31 % at lead 1 ($p = 0.199$) and nothing at longer leads. Stratifying by the share of each basin’s mascon-weighted signal that comes from land outside the basin gives what pooled effect there is a leakage signature: among basins large enough for GRACE to resolve it sits in the shared-footprint stratum (+1.16 %, $p = 0.0045$) and is negative elsewhere (−0.50 %); one small low-sharing cell resists this reading and is flagged in the text. Supplied instead as a state propagated to the forecast target time and applied as a residual correction on a deep sequence model, the same neighbor signal adds +0.9 to +2.0 % at every lead 1–6 (all $p \leq 3.2 \times 10^{-8}$), beats 20 of 20 degree-matched random-graph placebos in every lead–seed cell, and is positive in all 60 strata-by-lead cells with no concentration in shared-footprint basins – so the leakage explanation that fits the linear pocket does not fit it. The propagation, not the correction stage, is what makes it usable: given the neighbor’s raw 12-month history instead, the identical

correction stage recovers nothing. Short-lead TWSA forecasting is a filtering problem; beyond two to three months it becomes a forcing problem, and the two skill sources are complementary rather than competing.

1 Introduction

The GRACE and GRACE-FO satellite missions have measured monthly changes in terrestrial water storage (TWS) since 2002, integrating groundwater, soil moisture, surface water, snow, and ice into a single observable that no in situ network can replicate (Tapley et al., 2004; Landerer et al., 2020). Anomalies of this quantity (TWSA) carry months of memory, which makes them both a target and a source of predictability: TWSA forecasts support drought early warning and water management (Li et al., 2026; Arsenault et al., 2020), flood-potential assessment (Reager et al., 2014), and the bridging of the 2–3 month latency of GRACE-FO data releases (Mo et al., 2025).

A growing literature forecasts TWSA directly from data. Neural networks have been trained on GRACE and hydrometeorological series at basin scale (Ahmed et al., 2019), at grid scale globally (Li et al., 2020, 2024; Li and Kusche, 2026), for Europe with subsequent assimilation (Li et al., 2025), and for latency-filling (Mo et al., 2025); statistical time-series methods have been applied per basin (Ahi and Cekim, 2021); benchmark predictability has been quantified at decadal scale (Zhu et al., 2019); and physically based systems forecast TWS with seasonal meteorological forcing (Li et al., 2026; Getirana et al., 2020). Nearly all of these studies quantify their skill against some form of persistence or climatology, and the implicit logic is that whatever beats persistence has extracted predictive structure beyond simple memory.

This paper starts from an observation that undermines that logic for GRACE specifically. Basin-average TWSA is not a state variable observed cleanly; it is a state variable observed through a noisy instrument, with observation errors that are large for small and arid basins (Save et al., 2016; Landerer et al., 2020). A persistence forecast – last month’s observed anomaly, damped or not – propagates the observation noise of the anchor month into every forecast. The proper persistence-class reference for a noisy observable is not persistence of the observation but persistence of the *filtered state*: an estimate of the signal with the observation noise removed. For an AR(1) signal observed with white noise this optimal filter is the Kalman filter (Kalman, 1960), and it costs three parameters per basin. If this stronger member of the persistence class beats standard persistence by several percent, then part of the reported skill of learned models over persistence is not learned structure at all – it is noise filtering that a two-line state-space model performs for free.

The same benchmark-hygiene argument has been made in sea-ice forecasting, where Niraula and Goessling (2021) showed that damped-persistence references that “do not exploit observations optimally” flatter dynamical forecast systems, and built a stronger persistence-class benchmark from spatial damping. We execute the analogous move for satellite-observed TWSA via explicit signal/observation-noise separation. Within the GRACE literature, Kalman machinery is well established but serves different purposes: ensemble Kalman filters assimilate TWSA into land-surface models (Springer et al., 2026), and state-space smoothers densify or clean the gravity-field record itself (Kankanige et al., 2026). Neither uses a univariate AR(1)-plus-observation-noise model on the finished basin series *as a forecast benchmark*, which is the role it plays here; our filter is not a data-assimilation system and updates no model states. Related in spirit, Ahi and Cekim (2021) applied exponential-smoothing

family methods per basin without a noise-separation argument, and Kankanige et al. (2026) found persistence-based prediction of GRACE spherical-harmonic coefficients competitive at sub-monthly scales – independent support that the persistence class is underrated. Most directly, Nie et al. (2025) showed on 515 basins that deep architectures do not beat linear regression for TWS prediction when autoregressive information is deliberately excluded from the feature set; we show the complementary result that when autoregressive information is *included*, how one filters it dominates the model comparison.

If short-lead skill is mostly filtering, what is left for models to earn? The seasonal hydrologic forecasting literature answered a structurally identical question long ago: skill at short leads comes from initial conditions, and skill at long leads from meteorological forcing, with a crossover in between (Wood and Lettenmaier, 2008; Shukla and Lettenmaier, 2011; Koster et al., 2010; Yossef et al., 2013; Pechlivanidis et al., 2020). That partition was established with physically based models in ESP/reverse-ESP experiments, mostly for streamflow and soil moisture. It has not previously been located on satellite-observed TWSA, nor measured between two real, published, data-driven forecast systems. We do both, comparing our observation-noise-aware system head-to-head with the GRACE-FCast hindcasts of Li and Kusche (2026), whose LSTM ingests exogenous hydrometeorological predictors and, by design, no GRACE lags. The comparison locates an empirical crossing at approximately two months. We describe it as consistent with the initial-conditions/forcing partition rather than as a decomposition of skill into sources: the two systems differ along many axes at once, so what is measured is the lead profile, not the effect of any single design choice.

The third question the paper answers is whether TWSA fields themselves carry usable cross-basin information – whether one basin’s history helps forecast another’s future once own-basin filtering and local meteorology are accounted for. Spatial and graph-based TWSA methods increasingly assume so: Li and Kusche (2026) shift predictors across space, Arzoumanidis et al. (2026) build a proximity-plus-correlation basin graph for TWSA reconstruction, and Steidl and Zhu (2025) forecast TWSA with a hierarchical basin-graph network. None of these works isolates the incremental contribution of the cross-basin information, and none tests the graph against a null; Steidl and Zhu (2025) additionally target a reconstructed rather than observed TWSA product. To our knowledge, no prior work measures the incremental cross-basin information in *observed* GRACE TWSA against null models. We do so with degree-matched random graphs, a mascon leakage distance exclusion, a resolution stratification, and full ERA5 local-meteorology conditioning. The answer is two-sided, and the two sides must be kept apart: at the linear tier the incremental effect is a null, while the same neighbor signal does deliver skill through a dedicated nonlinear correction stage. We therefore also show what governs whether the signal is usable by a deep model at all, and it is not the architecture but the *representation*: the neighbor helps only when it arrives as a state already propagated to the forecast target time, and is inert as raw history however it is delivered.

The contributions, in order:

1. **An observation-noise-aware benchmark.** A per-basin three-parameter Kalman filter (AR(1) signal plus observation noise) beats damped persistence by 5.0–8.8% at leads 1–2 and per-basin ridge regression at five of six leads (DM $p \leq 0.013$; lead 4, $p = 0.078$, is not significant), making the filtering class the strongest own-basin reference at every lead. Consequence: benchmarked against persistence, a substantial share of published “ML skill” on deseasonalized TWSA is reproduced by three-parameter state-space filtering alone. We frame this constructively: the interesting question becomes

what *does* beat the filter, and the answer – exogenous forcing, a longer window of that forcing read linearly, and a small spatial correction – is a map of where genuine information lives.

2. **The initial-conditions/forcing crossover, located on observed TWSA.** Against the published GRACE-FCast hindcast (Li and Kusche, 2026; Li, 2025) on 227 matched basins under a common protocol, our best system wins lead 1 by 20.0 % (DM $p = 3.4 \times 10^{-4}$), is not detectably different at lead 2, and loses leads 3–6 by 13–30 %. The two systems are complementary, not competing: short leads are a filtering problem, long leads a forcing problem.

3. **Cross-basin information: a linear null, a nonlinear signal, and a representation that decides both.** The result is two-sided and we report both sides. *At the linear tier it is a null.* The registered capacity-matched contrast adds +0.31 % at lead 1 ($p = 0.199$) and is negative and non-significant at leads 2–6. What pooled lead-1 effect survives is a *leakage pocket*, not a size effect: splitting basins by how much of their mascon-weighted signal comes from land outside the basin concentrates the estimate in the shared-footprint stratum, at equal magnitude among basins GRACE resolves (+1.16 %, $p = 0.0045$) and among smaller ones (+1.17 %, $p = 0.0071$), while low-sharing resolved basins run *negative* (−0.50 %; one small low-sharing cell resists, Sect. 4.4). It is also short-horizon, dying after lead 2 (Sect. 4.4). *At the nonlinear tier it is real, placebo-validated, and does not carry that signature.* A neighbor-only residual stage on a deep own-state + ERA5 sequence model adds +0.91 to +1.96 % at every lead 1–6 (all $p \leq 3.2 \times 10^{-8}$), beats 20 of 20 degree-matched random-graph correction placebos in all twelve lead–seed cells, and is positive in all 60 strata-by-lead cells without concentrating in shared-footprint or small basins – it is worth +0.89 to +2.00 % in the very stratum where the linear effect is −0.50 %, and gains least in the small basins an artifact would favour. Meanwhile the same neighbor’s state fed to the same encoder as an input channel (there as a 12-month history; Sect. 4.5) yields no robust ensemble gain (+0.36 %, $p = 0.090$, at lead 1; negative at leads 2–6). A representation-matched arm identifies which difference is doing the work: the correction stage fed that same 12-month history instead of the propagated scalar also recovers nothing (not significant at any lead, and −0.87 to −2.24 % against the scalar correction at every lead), so it is the representation and not the delivery that separates the arms. The claim is therefore that cross-basin information exists, is invisible to standard linear methods and to raw-history deep-model inputs, and becomes usable once the neighbor is propagated to the forecast target time by its own estimated dynamics; we add a null-model control showing that what is found is the neighbor’s *information* rather than the two-stage architecture, and a stratification showing it is not the footprint-sharing artifact that explains the linear pocket.

Section 2 describes the data, Sect. 3 the decomposition, benchmark, controls, and learned models, Sect. 4 the results in the order above, and Sect. 5 the implications, the comparability of published numbers, and the limitations.

2 Data

2.1 GRACE/GRACE-FO mascons and basin sample

125 We use CSR RL0603M mascon solutions (Save et al., 2016) on a 0.25° grid, spanning April 2002 through May 2026, expressed
as equivalent water height (cm); the distributed file is titled RL0603M and reports product version RL06.2. Each of the 257
monthly solutions is assigned to its official calendar month using the product’s documented missing-month list, with the
solution’s data span asserted to overlap its assigned month. Midpoint binning mislabels the November 2011 and May 2015
130 solutions, whose arcs are centered on 31 October 2011 and 27 April 2015; under that convention two months hold two solutions
each and November 2011 and May 2015 are silently dropped from the record. This was a real defect in an earlier version of
our pipeline and is fixed here. Basins are defined by a HydroSHEDS-based Level-3 basin–mascon partition (Lehner and Grill,
2013); each basin series is the cosine-latitude area-weighted mean of its mask cells, with longitude centroids computed by
circular mean. From the 284 basins in the partition we exclude 46 ice-sheet basins (Antarctica, Greenland) and 4 water bodies,
leaving **234 basins** on all inhabited continents; 38 of these are African (37 mainland plus Madagascar), and 8 are glaciated and
135 tracked as a separate stratum.

2.1.1 Basin areas and native mascon geometry.

Basin areas used in the stratifications of Sect. 4.4 are our own cosine-latitude integration of the mask file. Native mascon
assignment comes from CSR’s official ancillary products: the RL06.3 native-mascon-ID mapping file, which labels every 0.25°
grid cell with its native mascon, and the RL06 v02 land/ocean masks (doi:10.15781/cgq9-nh24; Save et al., 2016), which CSR
140 recommends for basin definitions to minimize coastal leakage. As an internal check we also recovered the tiles empirically
– the gridded product replicates one tile value across the 0.25° cells it covers, so fingerprinting eight months and grouping
cells on exact series equality reconstructs the partition – and the two agree *exactly*: 42,107 tiles in both, cell-level purity 1.0
in both directions, adjusted Rand index 1.0, and per-basin contamination values identical to four decimals. The official tiles
have median area 12,123 km² (IQR 11,518–13,032), with cell counts rising from 16 to 24 with latitude – near-constant area at
145 varying cell count, CSR’s $\sim 1^\circ$ equal-area geodesic grid. A 90,000 km² basin therefore spans roughly *seven* native mascons,
not one; only 3 of the 234 basins span fewer than two. We accordingly do not describe that threshold as one resolution element.
CSR separately advises caution for basins below approximately 200,000 km², a cut under which 35 of our 199 nominally
resolved basins fall; Sect. 4.4 therefore repeats the footprint stratification at that threshold as a sensitivity.

2.2 ERA5-Land forcing

150 We use 11 ERA5-Land variables (Muñoz-Sabater et al., 2021) at 0.5° (server-side regridded from 0.1°), January 2001 through
May 2026: total precipitation, evaporation, total/surface/subsurface runoff, volumetric soil water in four layers, 2-m tempera-
ture, and snow depth water equivalent. Accumulated fluxes are converted from mean daily rates to monthly totals (cm month^{−1}),
evaporation is sign-flipped so positive means moisture loss, temperature is converted to $^\circ\text{C}$, and snow depth to cm water equiva-

lent. Basin aggregation uses the same cosine-latitude weighted means as the TWSA aggregation, mapping each 0.25° mask cell to its nearest 0.5° ERA5 cell. The resulting table ($284 \text{ basins} \times 305 \text{ months}$) is gapless. Because ERA5-Land is defined over land only, coastal and island basins aggregate their forcing from the land fraction of their mask: of the 234 analysis basins, 79 draw less than 90 % of their mask weight from valid ERA5 land cells and 13 less than 50 % (minimum 24.5 %). These basins’ forcing features average a sliver of coast and are correspondingly noisier; no basin is excluded for this, and the per-basin ERA5 results (Sect. 4.2) should be read with it in mind.

160 2.3 Climate indices

Twelve monthly climate indices (ENSO, IOD, NAO, and others) from NOAA PSL are used in the conditioning controls (Sect. 3.4) and in the pooled ridge’s climate-index variant (Table 2); AMM and PDO end before the test window and are dropped from that variant. No headline model takes index features.

2.4 The GRACE-FCast hindcast

165 For the head-to-head comparison we use the public CSR-FCast hindcast of Li and Kusche (2026), downloaded from PAN-GAEA (Li, 2025): monthly 1° global forecasts at leads 1–12 months, initialized December 2009 through May 2024 (174 initializations). Their LSTM forecasts TWSA from antecedent hydrometeorological and oceanographic predictors with correlation-based spatial predictor shifting; by design it takes no GRACE lags as input (Li and Kusche, 2026). We evaluate both their full-signal product and their non-seasonal variant (Sect. 3.6).

170 Two features of that choice bear on how our comparison should be read. GRACE-FCast is released as three products, built on the JPL, CSR and GSFC mascon solutions, and the skill figures reported in Li and Kusche (2026) are for the average of the three; the CSR member we use is the one their own training-residual diagnostic ranks least certain (3.0 versus 2.8 cm). We nonetheless compare against CSR-FCast alone because our own system is built on CSR mascons, so this is the mascon-matched contrast and the only one in which the two systems share a gravity solution — but it means our numbers characterize that member, not the product their headline scores describe. Second, their decomposition removes a linear trend as well as the seasonal cycle, so their non-seasonal variant and our deseasonalized target are defined on the same footing (Sect. 3.1); the functional form of their seasonal and trend fit follows Li et al. (2020) and is not restated in the forecast paper, which is why we also report the full-signal variant, where any such convention difference is absorbed by subtracting our own climatology.

3 Methods

180 3.1 Target definition: fold-safe deseasonalization

For each basin and each evaluation fold we fit, on the training window only, a climatology consisting of a linear trend plus annual and semiannual harmonics (six coefficients, ordinary least squares). The frozen coefficients deseasonalize both training and test data; the residual is divided by its training-window standard deviation. The forecast target is this deseasonalized,

per-basin-standardized residual – the interannual anomalies that drought and flood applications need – and all pooled metrics
185 are computed in this standardized space so that high-variance basins do not dominate. Physical-unit (cm) results are reported
where noted.

3.2 Evaluation protocol

3.2.1 Folds.

Five expanding-window folds whose issue windows jointly cover June 2019 through April 2026 at lead 1, with targets reaching
190 May 2026 (17, 17, 17, 17, and 15 issue months per basin at lead 1, one fewer per basin at each successive lead as the shifted
target runs off the end of the record). The matched sample is therefore $n = 19,422/19,188/18,954/18,720/18,486/18,252$
rows at leads 1–6, identical for every model at a given lead. Fold membership is defined on forecast *issue* dates: within
each fold, every climatology, filter parameter, model coefficient, neighbor graph, and standardization moment is fit strictly on
observations before the fold’s first issue date, and every test forecast is issued at or after that date – the model is frozen once
195 and then used forward, as it would be in deployment. (An earlier design split on target dates, which let lead-2–6 transforms see
up to $h - 1$ months past the issue date; a 2026-08-13 audit caught this and the entire pipeline was rerun under the issue-date
protocol.) Earlier issue windows are unusable: the 2017–2018 gap between GRACE and GRACE-FO destroys the lag windows
needed to forecast across them. All results are retrospective potential skill: issue time is defined relative to the last available
observation, so the 2–3 month GRACE-FO release latency and ERA5 availability at issue time are not modeled.

200 3.2.2 Metrics and tests.

The primary metric is pooled RMSE in standardized units, converted to a skill score $1 - \text{MSE}_A/\text{MSE}_B$ against a stated
reference. Pooled model comparisons use the Diebold–Mariano test (Diebold and Mariano, 1995) with the small-sample cor-
rection of Harvey et al. (1997) on cross-basin monthly mean loss differentials, with Newey–West HAC variance (Newey and
West, 1987) (maximum lag $\max(h - 1, 1)$); confidence intervals use a moving-block bootstrap (Künsch, 1989) (block length
205 $\max(3, h)$, 2000 draws, so the block covers the lag- $(h-1)$ dependence of overlapping h -step forecasts). Per-basin signifi-
cance uses per-basin DM statistics on the scored test months – 83 at lead 1 declining monthly to 78 at lead 6, the last issue
month receding from April 2026 to November 2025 as the shifted target runs off the record – with Benjamini–Hochberg false-
discovery-rate control at $q = 0.10$ (Benjamini and Hochberg, 1995), following the field-significance reasoning of Wilks (2016);
we always report helped and hurt basins separately.

210 3.3 Baseline ladder and the observation-noise-aware benchmark

The baseline ladder comprises: zero-anomaly climatology; persistence (last observed anomaly); damped persistence in two
variants (an AR(1) damping coefficient fit per basin, and a regression-estimated variant), of which we always score against the
stronger variant at each lead; a pooled ridge on own-basin lagged anomalies ($\alpha = 1$, lags 0, 1, 2, 5, 11 months); the same ridge
with climate-index lags; and a per-basin ridge.

215 The benchmark this paper argues for models each basin’s deseasonalized residual y_t as an AR(1) latent signal x_t observed with additive white noise:

$$x_t = \rho x_{t-1} + \epsilon_t, \quad \epsilon_t \sim \mathcal{N}(0, q), \quad (1)$$

$$y_t = x_t + \eta_t, \quad \eta_t \sim \mathcal{N}(0, r), \quad (2)$$

with (ρ, q, r) estimated per basin by maximum likelihood (L-BFGS-B, multiple starts) on pre-test data. The h -month forecast
 220 is $\hat{y}_{t+h|t} = \rho^h x_{t|t}$, damped persistence of the *filtered state*. The point forecast depends on ρ and the noise ratio q/r only; the filter handles the mission gap and missing months by skipping update steps, which inflates state uncertainty without biasing the state. We call this forecast the *Kalman forecast*. Two further members of the same “filtering class” reuse the filtered state: the *ridge on filtered states*, a pooled ridge on lagged filtered states, and the neighbor-augmented variants of Sect. 3.4. We emphasize what this benchmark is not: it assimilates nothing into any model and is unrelated to the GRACE data-assimilation
 225 literature reviewed by Springer et al. (2026); it is a forecast reference with three scalars per basin.

3.3.1 Identification caveat, and the $r = 0$ ablation.

The decomposition into signal variance q and observation-noise variance r is a modeling choice, not a measurement: high-frequency hydrology, non-AR(1) dynamics, and decomposition error are absorbed into r alongside true instrument noise, and in a substantial minority of basin–fold fits the likelihood pushes r to its lower boundary: all 1170 basin–fold fits converge, but
 230 259 of them land at $r < 10^{-6}$, and 31 basins sit at the boundary in all five folds. So r is not, per basin, a clean instrument-noise estimate. Whether the benchmark’s *margin* is nonetheless the filtering term is a separate, answerable question, and an ablation answers it: refitting the identical model class by the identical MLE with r fixed at zero – which forces the update gain to one, so the “filter” tracks the raw observation and the model degenerates to damped persistence with MLE ρ – removes the entire margin and more. The free- r filter beats its $r = 0$ twin by +5.4% at lead 1 rising to +12.7% at lead 5 (DM $p < 10^{-13}$ at every
 235 lead), while the $r = 0$ twin *trails* the tuned damped-persistence baseline at leads 2–6 (−0.9 to −11.6%): the AR(1) structure alone, without the noise term, is worth nothing beyond persistence at any lead past the first. What is filtered out doubtless mixes instrument noise with fast unmodelled hydrology; the ablation attributes the margin to the act of filtering the observation before extrapolating, not to a specific noise source.

3.4 Neighbor graphs, controls, and the jump screen

240 3.4.1 Graphs.

For each target basin and fold we select one source basin (top-1) by: training-window Pearson correlation of the TWSA series (*correlation graph*, the headline treatment); correlation of the source’s one-month-lagged series with the target (*predictive-lag graph*); or great-circle proximity (*geographic graph*). Neighbor information enters as the lagged, forward-propagated filtered state of the source basin appended to the target’s feature set. Distance-restricted variants exclude sources within 300, 500, or
 245 1000 km.

3.4.2 Degree-matched random graphs.

This is the primary control throughout the paper. Every neighbor model is compared against 20 or 50 placebo graphs that preserve in-degree but draw source basins uniformly at random. The placebo models are refit end-to-end, so the comparison is capacity-matched; beating all N draws gives an empirical rank $p_{\text{rank}} = 1/(N+1)$. Two properties of this statistic must be kept in view. First, p_{rank} is bounded below by $1/(N+1)$, so at $N = 20$ every “significant” cell is pinned at 0.0476 and the rank p understates the evidence; where the margin matters we therefore also report the real graph’s position in placebo standard deviations. Second, the placebo draws are redrawn independently for each fold and horizon (an earlier version of this analysis reused one draw set across cells; a 2026-08-15 repair removed that), but seeds within a cell share the cell’s draws, so per-seed rank results at a given lead are re-scorings of the same 20 graphs (Sect. 5.5).

3.4.3 IAAFT surrogates.

Iterated amplitude-adjusted Fourier transform surrogates (Schreiber and Schmitz, 1996, 2000) preserve each basin’s marginal distribution and power spectrum while destroying cross-basin phase alignment. The transform runs on the full monthly grid with GRACE gap months linearly interpolated and the gap pattern re-imposed afterward, so the preserved spectrum belongs to the true monthly time axis rather than a gap-compressed one. The *real* selected graph is reused for all 99 surrogate datasets: the tested null is that the selected neighbor’s time alignment carries no information, with selection itself covered by the random-graph placebos. Surrogates are generated from the full record, test months included: no information flows into the real arm, and the null draws share the real series’ full-record statistics, which keeps the null closely matched to the real arm. Because the surrogate distributions are estimated with test-period observations, however, this is a transductive sensitivity rather than a fold-pure prospective control, and we do not claim the design is conservative.

3.4.4 Conditioning.

Climate-index conditioning gives ENSO and IOD lags to every model in a comparison; ERA5 conditioning gives all 11 local-meteorology variables (lags 0–2, deseasonalized, train-standardized, winsorized at $\pm 10\sigma$) to every model, while placebos re-randomize only the neighbor. The test is one-directional: if conditioning leaves the neighbor estimate unmoved, the neighbor is not acting as a proxy for shared local weather. It does not, on its own, establish that the neighbor carries a signal. The winsorization guard exists because near-zero training variance in arid basins otherwise detonates standardized ERA5 features (the 2023 Libya flood reaches several orders of magnitude beyond the training standard deviation).

3.4.5 Footprint-sharing stratification.

The physical leakage quantity a neighbor effect could be an artifact of is not basin size but the degree to which a basin’s mascon footprint reaches beyond the basin. We therefore define, per basin, a *contamination* score: the share of the basin’s mascon-weighted signal contributed by land outside the basin, $\sum_t (w_t/W)(1 - f_t)$, where w_t is the basin’s weight on native tile t and f_t is that tile’s fill fraction from within the basin. Contamination and area are nearly independent axes (Spearman

−0.325), so stratifying on one is not a restatement of the other. We report contamination terciles, the 90,000 km² area split retained from earlier drafts, and the 2 × 2 cross of the two. Two properties of the metric bound its interpretation: it treats the union of all 284 mask units as “land”, so the ocean half of a coastal tile is not counted as foreign and island and coastal basins score low; it is therefore a *lower bound* on footprint sharing.

3.4.6 Jump screen.

A 6 σ outlier screen flags basins with step-like jumps and removes them before the neighbor comparison, in a train-only version and in a test-informed variant kept as a post-hoc diagnostic only. On the corrected pipeline the train-only screen flags 2 basins (Taiwan; Northern Kordofan) and leaves the lead-1 neighbor comparison at +0.21 % against its own-ridge reference, still ahead of all 50 placebo draws ($p_{\text{rank}} = 0.02$); the test-informed variant flags 27 and gives +0.08 %, likewise 50/50.

3.5 Learned models

All learned models predict the residual of the Kalman forecast (the final prediction is the Kalman forecast plus the learned correction), and each is paired with a ridge “twin” on identical features so that architecture and information content are never conflated. Unless stated otherwise, deep models are run with multiple seeds and reported as seed ensembles.

3.5.1 Nonlinear heads.

Gradient-boosted trees and multilayer perceptrons on the same feature sets as the ridges.

3.5.2 Residual MLP.

Stage one: the ridge twin on own-basin features. Stage two: a small MLP (hidden layers 64–32, early stopping) trained on the stage-1 *training* residual using only the neighbor’s propagated filtered-state lags. Three seeds.

3.5.3 LSTM.

A shared-encoder LSTM (Hochreiter and Schmidhuber, 1997) over 12-month windows of the own filtered state concatenated with the 11 ERA5 channels, with a time-ordered ~85/15 train/validation split for early stopping; two seeds. Variants add the neighbor state as a thirteenth input channel.

3.5.4 The stacked system: LSTM plus a residual-correction stage.

Stage one: the own-state + ERA5 LSTM above, no neighbor. Stage two: an MLP trained on the stage-1 training residual from the propagated neighbor state only. Placebo graphs re-randomize only the stage-2 neighbor and ride the stage-1 network of the matching seed, so the placebo comparison is like-for-like within each seed. This construction matters: scoring one seed’s model against another seed’s placebos introduces a handicap several times larger than the placebo spread itself. The same training run also produces a neighbor-free *own-state control*, which sends each basin’s *own* propagated state through the identical stage-2

305 path – same stage-1 network object, same stage-2 model, capacity, and seed – so that only the information source differs. We report it as a diagnostic rather than as the primary control, for reasons given in Sect. 4.5.

3.5.5 Graph attention network.

A one-layer GAT (Veličković et al., 2018) with one- and two-hop variants, three seeds, included to test learned message passing.

Table 1 collects the model names used in the rest of the paper; each name is used for exactly one model throughout.

310 3.6 Comparison protocol for GRACE-FCast

The Li (2025) hindcast is aggregated to our basin geometry by exact $0.25^\circ \rightarrow 1^\circ$ cell-containment weights; 7 island basins with zero coverage in their land mask drop out, leaving **227 matched basins**. Their forecasts are moved into each fold’s deseasonalized standardized space using our train-window climatology plus a per-(fold, basin, lead) mean offset estimated on pre-test data only. The matched sample is the identical set of (basin, target-month) rows for every model at every lead: 60 test
315 months, folds 1–4; fold 5 lies entirely beyond their final initialization. This gives a complete rectangle of $227 \times 60 = 13,620$ rows per model at each lead. An independent audit reproduced the aggregation, alignment, and every headline number digit-for-digit (maximum difference 0.0 across 30 recomputed cells).

Two protocol asymmetries favor their system and are carried as caveats wherever it wins: their training uses the Yin et al. (2023) gap reconstruction (hindsight information across the 2017–2018 gap), and their product re-extrapolates seasonal and
320 trend components monthly whereas our climatology is fold-frozen – a difference worth about 0.07–0.09 in skill units at leads 4–6, bounded by their non-seasonal variant, which still wins those leads.

4 Results

4.1 The benchmark: persistence of the filtered state

Table 2 and Fig. 1 show the matched-row baseline ladder. The filtering class supplies the strongest own-basin reference at
325 every lead. The bare Kalman forecast beats the stronger damped-persistence variant by +4.98%, +8.79%, +5.62%, +3.07%, +2.55%, and +3.63% at leads 1–6 ($DM\ p \leq 8 \times 10^{-7}$ everywhere). It beats the per-basin ridge at five of six leads (+1.4 to +3.9%, $DM\ p \leq 0.013$); at lead 4 the margin is +1.04% and the DM test does not reject ($p = 0.078$), so we do not count that lead. (The moving-block bootstrap interval at lead 4 excludes zero; the DM test is the criterion we registered and the one we report.) Against the *pooled* ridge the filter is significant at leads 1–2 only (+3.86% and +4.21%); at leads 3–6 the two are
330 not separable, and at leads 4–5 the filter is nominally behind (−0.59%, −0.64%, both n.s.). At leads 3–6 the filtering class is instead carried by a pooled ridge on lagged *filtered* states, which beats the per-basin ridge by +3.4 to +5.4% at leads 4–6 ($p \leq 2.1 \times 10^{-6}$); the signal/noise separation, not the ridge head, is what carries the class.

Three readings of Table 2 matter for the field. First, the gap between raw persistence and damped persistence (16–23%) is the familiar cost of not damping; the further 5.0–8.8% from damping the filtered state instead of the observation is the

Table 1. Model names used throughout the paper, with the defining subsection. Appendix A maps each name to its identifier in the archived result files.

Name	Definition
Climatology	Zero-anomaly forecast (Sect. 3.3).
Persistence	Last observed anomaly (Sect. 3.3).
Damped persistence	Last anomaly damped toward zero; two per-basin damping variants, always scored against the stronger at each lead (Sect. 3.3).
Pooled ridge	Ridge on own-basin lagged anomalies, one model for all basins; a variant adds climate-index lags (Sect. 3.3).
Per-basin ridge	The pooled ridge’s features, refit separately for each basin (Sect. 3.3).
Kalman forecast	The proposed benchmark: damped persistence of the filtered state from the per-basin AR(1)-plus-observation-noise model (Sect. 3.3).
Ridge on filtered states	Pooled ridge on lagged filtered states (Sect. 3.3).
Neighbor ridge	Ridge on filtered states plus one selected neighbor’s propagated filtered state; graph-construction, distance-restricted, and conditioned variants (Sect. 3.4).
Own-basin ridge / ERA5 ridge	Ridge twins on the basin’s own features, without and with the ERA5-Land lag features (Sect. 3.5).
GBM head / MLP head	Gradient-boosted trees and multilayer perceptron on the same features as the corresponding ridge twin (Sect. 3.5).
Residual MLP	Two-stage model: a ridge first stage plus a small MLP trained on the stage-1 training residual; variants differ in the features each stage receives (Sect. 3.5).
LSTM	Shared-encoder LSTM over 12-month windows of the filtered state and the ERA5-Land channels; the neighbor-channel variant adds the neighbor state as a thirteenth input channel (Sect. 3.5).
Stacked system	The LSTM as stage 1 plus a neighbor-only MLP correction trained on the stage-1 training residual (Sect. 3.5).
Own-state control	The stacked system with stage 2 fed the basin’s own propagated state instead of a neighbor’s; a diagnostic, never a skill number (Sects. 3.5 and 4.5).
GAT	One-layer graph attention network, one- and two-hop variants (Sect. 3.5).
GRACE-FCast	The published forecast system of Li and Kusche (2026); full-signal and non-seasonal products (Sects. 2.4 and 3.6).

335 filtering-class margin, and the $r = 0$ ablation of Sect. 3.3 identifies it as the noise-filtering term specifically: the same model class with the noise variance removed gives the margin back and more, trailing damped persistence itself at leads 2–6. The margin is of the same order as many published model-over-persistence margins. without noise separation sit *between* the two persistence variants and the filter: the per-basin ridge – five lag coefficients per basin – never reaches the three-parameter filter

Table 2. Pooled skill (%) versus the stronger damped-persistence variant at each lead, matched rows across all models (234 basins, five folds; $n = 19,422/19,188/18,954/18,720/18,486/18,252$ at leads 1–6, constant across models; see Sect. 3.2). Positive is better. The reference is ρ -damping at lead 1 and regression damping at leads 2–6, whichever is stronger; the ρ -damped variant alone degrades to -6.4 to -12.0% at leads 3–6 and scoring against it would flatter every other row. The Kalman forecast is damped persistence applied to the filtered state rather than to the raw observation; “ridge on filtered states” is a pooled ridge on lagged filtered states. DM p values for the two filtering-class rows are $\leq 8 \times 10^{-7}$ at every lead.

Model	$h=1$	$h=2$	$h=3$	$h=4$	$h=5$	$h=6$
Climatology (zero anomaly)	−125.6	−70.7	−58.9	−51.4	−42.7	−33.6
Persistence	−20.7	−22.7	−21.4	−19.5	−16.0	−15.8
Damped persistence (stronger variant)	0.0	0.0	0.0	0.0	0.0	0.0
Pooled ridge (own lags)	+1.2	+4.8	+5.1	+3.6	+3.2	+2.5
Pooled ridge + climate indices	+1.0	+4.5	+5.1	+3.7	+3.2	+2.4
Per-basin ridge (own lags)	+3.1	+5.9	+4.3	+2.1	+0.4	−0.3
Kalman (filtered-state persistence)	+ 5.0	+ 8.8	+5.6	+3.1	+2.5	+3.6
Ridge on filtered states	+4.8	+7.9	+ 6.0	+ 5.4	+ 5.1	+ 5.2

Skills are recomputed on matched rows from the archived per-prediction files; the archived source file for every table and figure is listed in Appendix A.

at any lead, and by lead 6 it has fallen below damped persistence altogether. Third, the benchmark is nearly free: it adds no data requirements and two lines of state-space code, so we propose it as the default persistence-class reference for deseasonalized GRACE basin forecasting. The reading is constructive: the ladder does not say learned models are useless; it says the bar they must clear sits higher than the field’s standard reference, and the rest of this section is a map of what clears it.

4.1.1 Robustness rather than growth.

These numbers were produced after two evaluation defects were found and fixed in our own pipeline: two CSR solutions had been binned into the wrong calendar month, and fold membership had been defined on target dates rather than forecast issue dates (Sect. 3.2). Both fixes tighten the protocol, and neither moves the benchmark result: the filter’s margin over damped persistence was $+4.8/+8.1/+5.3/+3.2/+2.5/+3.5\%$ before the corrections and $+5.0/+8.8/+5.6/+3.1/+2.5/+3.6\%$ after them. We report this as a robustness property of the benchmark, not as an improvement: the margin did not grow. What the corrections did change is the significance of the lead-4 contrast against the per-basin ridge and the reading of the pooled ridge, both of which are stated above in their corrected form.




Figure 1. The benchmark ladder as skill versus lead. Pooled skill (%) relative to the stronger damped-persistence variant at each lead for the models of Table 2, with block-bootstrap 95 % confidence intervals on the two filtering-class curves. The gap between the damped persistence of the raw observation (zero line) and damped persistence of the Kalman-filtered state is the gain from filtering the anchor observation: 5.0–8.8 % at leads 1–2. Conventional statistical models (pooled and per-basin ridge) sit between the two persistence classes at every lead.

4.2 What beats the filter: forcing, and a longer window of forcing

4.2.1 Exogenous forcing is the largest single skill source, and it is a lead-1 effect.

Adding ERA5-Land lag features to the own-basin ridge yields +4.62% at lead 1 (95 % CI +1.9 to +6.8, DM $p = 1.0 \times 10^{-4}$) – an order of magnitude larger than anything the neighbor channel of Sect. 4.4 produces at the linear tier. Beyond lead 1 it is not detectable: +1.39% at lead 2 ($p = 0.092$) and +0.33% at lead 3 ($p = 0.61$). Forcing supplied to a linear head therefore buys one month; carrying it further requires the longer forcing window below.

Three decompositions locate the gain. By variable (leave-one-out and solo refits of the same ridge arm), the information is spread across the collinear moisture variables rather than owned by any one of them: evaporation carries the largest unique contribution (+1.45 percentage points; drop-test $p = 0.002$, the largest of the eleven leave-one-out tests and the only one that survives Bonferroni correction across that family with room to spare), the top soil-moisture layers are the strongest standalone predictors (+2.3 and +2.2% alone, $p < 10^{-3}$), and no other variable contributes as much as 0.6 of a point uniquely. Snow adds nothing, and two variables – 2-m temperature and the top soil-moisture layer – have small *negative* unique terms (dropping either slightly improves the arm, $p \leq 0.04$), as densely collinear features can. By fold, the gain is robustly positive in four of

five test windows (+6.0 to +8.5%, $p \leq 0.002$) and fails in exactly one – the May 2022–September 2023 window spanning the triple-dip La Niña (−3.0%, $p = 0.32$) – the one climate state in the record where reanalysis forcing misinformed the model. And by geography, the gain is a northern-hemisphere phenomenon: pooled skill is +7.1, +6.9, and +6.0% in North America, Europe, and Asia, +3.4% in the Australia–Oceania group ($p = 0.053$), and statistically zero in South America (+1.5%, $p = 0.66$) and Africa (−0.4%, $p = 0.87$). Per basin, 23 of 234 survive FDR control (21 helped, 2 hurt). The mirror-image relation to the neighbor channel that Sect. 5.3 builds on – ERA5 weakest exactly where the neighbor signal is strongest – rests on these corrected numbers.

4.2.2 A longer window of forcing extends the gain; sequence architecture does not.

On identical own-basin features, gradient boosting and MLP heads never beat their ridge twins; across every feature set the sole nominal exception is the gradient-boosted head on the neighbor-augmented features at lead 2 (+1.12%, $p = 0.019$), which no other lead or head supports. With ERA5 aboard, the GBM is statistically tied at leads 1–2 (+0.14%, $p = 0.77$; +0.60%, $p = 0.13$) and significantly *worse* at lead 3 (−0.71%, $p = 0.037$), while the MLP is tied at lead 1 (−0.53%, $p = 0.30$) and significantly worse at leads 2–3 (−2.56% and −3.14%, $p \leq 7.2 \times 10^{-5}$). This is the within-study analogue of Nie et al. (2025), and the corrected pipeline puts it on firmer ground than our earlier evaluation did. The apparent exception was temporal, and dissecting it changed our reading of it. An LSTM over 12-month windows of filtered state plus ERA5 channels beats the ridge by +2.17%, +2.36%, and +1.43% at leads 1–3 (two-seed ensemble, DM $p \leq 4.6 \times 10^{-5}$) – but that ridge sees only issue-time lag features, so the pair confounds the architecture with the amount of history it reads. Equalizing the information settles the attribution. A ridge over the identical 12-month history, flattened to a plain feature vector, gains +3.34, +5.01, and +3.49% over the lag-feature ridge – more than the LSTM’s entire margin – and head-to-head the LSTM *loses* to it at every lead (−1.20%, $p = 0.037$; −2.80%, $p = 1.8 \times 10^{-10}$; −2.13%, $p = 3.8 \times 10^{-9}$). One mechanical asymmetry in that pair deserves its own control: the LSTM holds out the last 15% of training months for early stopping, so it fits its weights on fewer rows than the ridge. Refitting the flat ridge on the LSTM’s exact 85% training rows changes it by less than 0.2% and leaves the comparison intact (+1.02, +2.76, +2.27% over the LSTM ensemble; lead 1 weakens to $p = 0.094$, leads 2–3 remain $p < 10^{-6}$), so the flat ridge’s advantage does not rest on the extra training rows. An MLP on the same flattened history loses to the flat ridge as well (−3.08 to −2.26%). What we had been calling a sequence-modeling gain is a history-length gain: twelve months of forcing carries real information, the linear reader extracts the most of it, and the sequence architecture pays a premium for re-learning what the flattening gives away for free. Without ERA5, the LSTM does not durably beat the ridge: +0.32% at lead 1 ($p = 0.10$), +0.93% at lead 2 ($p = 1.5 \times 10^{-4}$), and significantly *worse* at lead 3 (−1.03%, $p = 0.008$) – the filter has already extracted most of the temporal structure of the TWSA series itself, and the usable history is the forcing’s, not the TWSA’s. A one-layer graph attention network never beats its ridge twin at any seed, lead, or graph variant (−0.90, −0.24, −0.61% at leads 1–3 against its own-graph twin), and its two-hop neighborhoods are never better than one-hop. An exhaustive scan of the GAT variants turns up exactly one nominal win (+0.65% at lead 2, $p = 0.035$, ensemble only), which no individual seed supports and whose placebo ranks are poor; we read it as noise. Learned message passing at $n=234$ basins with ~65 training months per fold costs more than it captures.

Table 3. Head-to-head skill (%) of our stacked system (two-seed ensemble) versus the GRACE-FCast full product (Li and Kusche, 2026; Li, 2025) on the matched sample (227 basins, 60 months, 13,620 rows per lead). Positive means our system is better. Per-basin wins: basins with lower MSE for our system, of 227.

Lead	Skill (%)	95 % CI	DM p	Per-basin wins
1	+20.0	[+8.1, +30.5]	3.4×10^{-4}	137
2	-3.5	[-16.5, +7.6]	0.50	82
3	-12.9	[-24.0, -3.5]	0.015	71
4	-17.4	[-30.2, -8.3]	0.0021	65
5	-23.6	[-39.0, -15.3]	1.3×10^{-4}	59
6	-30.3	[-46.2, -22.8]	1.5×10^{-6}	59

Archived source files are listed in Appendix A.

4.3 The crossing: two published systems on observed TWSA

Table 3 and Fig. 2 report the matched-protocol comparison between our best system (the stacked system, Sect. 4.5) and GRACE-FCast (Li and Kusche, 2026).

The pattern is an empirical crossing consistent with the classic initial-conditions-versus-forcing crossover (Wood and Lettenmaier, 2008; Shukla and Lettenmaier, 2011; Yossef et al., 2013), here measured for the first time on satellite-observed TWSA between two real, published, data-driven systems. Our system wins lead 1 by +20.0%; the systems are not detectably different at lead 2 (-3.5%, $p = 0.50$, with an interval spanning -16.5 to +7.6% – an absence of evidence, not a demonstrated tie); theirs wins leads 3–6 with a gap that widens monotonically to -30.3%. We say “consistent with” rather than claiming a causal source decomposition: the two systems differ in architecture, training data, gap handling, seasonal treatment, and predictor sets simultaneously (and our own stack carries ERA5 soil-state information), so the lead profile, not any single design choice, is what is measured. The structural reading is nonetheless natural. Their LSTM takes no GRACE lags: at lead 1 the most recent observation is the single strongest predictor, and the bare Kalman forecast alone – three parameters, no exogenous data – beats their full product by +14.6% ($p = 0.0039$) there. The sharpest form of the split is against damped persistence on the matched sample: our system scores +11.1 to +16.6% at every lead 1–6, whereas theirs is nominally *worse* than damped persistence at lead 1 (-11.1%, $p = 0.057$ – suggestive rather than established) before rising to +14.8, +23.3, +29.0, +31.7, and +34.1% at leads 2–6. That lead-1 deficit is the weakest claim in the comparison and we do not lean on it: besides the p -value, it is a pooled-magnitude effect rather than a broad one, since their product still beats damped persistence in 115 of the 227 matched basins. Two framing points also belong with it. Damped persistence at lead 1 is not a baseline available to their intended user: mascon releases carry a 2–3 month latency, which is the stated motivation for their product, so a forecast conditioned on last month’s storage assumes an observation that user does not yet have. And Li and Kusche (2026) never benchmark against persistence, climatology, or any naive reference at any lead, so this is not a result their paper contradicts — it is a floor they did not test against. The comparison remains meaningful as a scientific skill floor, and the crossing itself does not rest on the




Figure 2. The initial-conditions/forcing crossover on observed TWSA. Skill (%) of our stacked system relative to the GRACE-FCast full product (Li and Kusche, 2026) as a function of lead on the matched sample (227 basins, 60 months), with block-bootstrap 95 % bands; a vertical line marks the crossover between leads 2 and 3. Secondary curves show the bare Kalman forecast against their product (structural lead-1 advantage of the filtered initial condition) and our system against their non-seasonal variant (the crossing survives removal of their seasonal/trend re-extrapolation edge). Diebold–Mariano significance is marked per lead.

420 lead-1 sign. Their exogenous-forcing design captures precipitation-driven variability at leads where the initial-condition signal has decayed; by lead 6 their skill over damped persistence is 0.34 versus 0.14 for our best system. The crossing is robust to their known protocol advantages: the non-seasonal variant – which removes the monthly re-extrapolated seasonal/trend edge – still beats our system at leads 4–6 (−8.9 to −23.4%, $p \leq 2.9 \times 10^{-3}$), while we beat it at leads 1–2 (+30.9%, +9.0%) and lead 3 is a wash (−1.7%, n.s.) – against that variant the crossing sits between leads 3 and 4 rather than between leads 2 and 3.

425 It is worth noting that Li and Kusche (2026) document the same target-definition sensitivity from the other side, and quantify it. Scoring their own product at lead 1, they report correlation above 0.8 over approximately 72% of global land for the full signal but only 20% once the seasonal cycle and linear trend are removed, and they attribute the difference to the extrapolated components directly: including them “improves correlation while increasing the overall error”. Most of the correlation in a full-signal GRACE forecast score is therefore carried by terms that are extrapolated rather than predicted. That is our benchmark
 430 argument restated in their numbers, and it is why the deseasonalized target is the one on which the two systems can be honestly compared at all.

Because no other published study reports our pooled deseasonalized metric, we also state physical units on the matched sample, both pooled and as a per-basin median – and the two tell usefully different stories. Pooled, our system’s RMSE is 5.0, 5.6, and 5.8 cm at leads 1–3 against their nearly lead-flat 7.6–7.7 cm; that comparison flatters us, because pooled cm error is dominated by a few high-amplitude Arctic glaciated basins where their product is weakest. The per-basin median is the honest form: 2.4 versus 2.7 cm at lead 1 in our favor, then 3.1 versus 2.8 and 3.3 versus 2.9 cm in theirs at leads 2–3. In median physical units the handover therefore arrives one lead earlier than in the standardized metric – the typical basin is better served by their product from lead 2 – which sharpens rather than weakens the complementarity reading: the filtering advantage is concentrated at lead 1 and in exactly the high-amplitude basins the pooled number overweights. The equal-weight standardized comparison of Table 3 remains our primary metric.

4.3.1 A descriptive best-of-both splice.

Splicing our forecasts at short leads onto theirs at long leads is a way to quantify the complementarity, and we report it as a descriptive sketch rather than a validated forecast product: the splice points would be chosen on the same test sample, several variants were examined without multiplicity control, and pooled DM differentials over all leads are diluted by the shared long-lead segment. Row-pooled over leads 1–6 on the matched sample, our system alone gains +11.6% against damped persistence and their product alone +23.8%, while the splices reach +24.4% (filtered-state short leads, $p = 3.0 \times 10^{-8}$) and +24.8% (ERA5-conditioned short leads, $p = 1.8 \times 10^{-9}$); restricting the handover to lead 1 alone gives +25.4 and +25.7%. The dilution the caveat above anticipates is visible in exactly those numbers: the pooled splice clears their product by well under a percentage point, because four of the six leads it pools over *are* their product. The gain lives where the splice actually differs from them — at lead 1, where it improves on their forecast by +14.5% (CI +3.8 to +24.4, $p = 0.004$) using our filtered state, or +17.3% (CI +4.6 to +28.4, $p = 0.003$) using the ERA5-conditioned variant. The per-lead tests in Table 3 are the evidence either way; the splice only illustrates what a filtering-plus-forcing hybrid could deliver.

4.4 Cross-basin information at the linear tier: a null result with a footprint-sharing pocket

At the linear tier the cross-basin channel does not pay. Adding one correlation-selected neighbor’s lagged filtered state to the ridge on filtered states – the registered, capacity-matched paired comparison, in which both models share the same head and own-basin features – changes pooled lead-1 skill by +0.31% (95% CI –0.09 to +0.82, DM $p = 0.199$) – not significant. At leads 2–6 the estimate is negative at every lead (–0.10, –0.19, –0.36, –0.35, –0.32%), also without significance; the nearest thing to a significant result anywhere in this family runs *against* the neighbor (lead 5, $p = 0.070$). We report this as a controlled negative result, and we are explicit that it reverses our own earlier finding: before the two evaluation defects of Sect. 3.2 were fixed, the same contrast read +0.49% with $p = 0.022$. What makes the null informative rather than merely inconclusive is the control battery the estimate was put through (Table 4 and Fig. 3), which no prior spatial-TWSA claim has faced, and which localizes the little that is there.

4.4.1 Two questions that must be kept apart.

The placebo evidence and the paired-contrast evidence answer different questions, and blurring them is how a null becomes an apparent finding. *Does correlation-selection beat a random neighbor?* At lead 1, yes: the real graph beats all 50 degree-matched random graphs ($p_{\text{rank}} = 0.0196$), and so does the ≥ 300 km-restricted variant. The selection step is doing something real. *Does adding a neighbor beat not adding one?* No ($p = 0.199$). The first question is about graph construction; only the second is about skill, and only the second is the paper’s claim. Beyond lead 1 both answers turn negative together: at leads 2–6 the real graph is beaten by 0 of 50 placebos, meaning the correlation-selected neighbor is worse than a randomly chosen basin. (The ≥ 300 km variant reaches 2 of 50 at lead 2; 0 of 50 elsewhere.)

4.4.2 Stratification (primary sensitivity): the effect is a footprint-sharing pocket, not a size effect.

Splitting the 234 basins at 90,000 km² puts the whole lead-1 estimate in the small stratum (35 basins, +1.033%, $p = 0.00065$) and none of it in the large one (199 basins, +0.099%, $p = 0.715$), and that split is not a knife edge: every cut from 60,000 to 200,000 km² gives the small stratum +0.9 to +1.3% and leaves the complement non-significant. It is tempting to read this as GRACE’s resolution limit, and an earlier version of this paper did. That reading is wrong.

Basin size is a proxy. The physical quantity is *footprint sharing* – how much of a basin’s mascon-weighted signal comes from land outside the basin (Sect. 3.4) – and it is nearly independent of area (Spearman -0.325). Crossing the two axes separates them:

Area stratum	Footprint sharing	Basins	Lead-1 skill (%)	DM p
Resolved ($\geq 90,000$ km ²)	high (top tercile)	57	+1.161	0.0045
Resolved ($\geq 90,000$ km ²)	low/mid	142	−0.500	0.082
Small ($< 90,000$ km ²)	high (top tercile)	21	+1.166	0.0071
Small ($< 90,000$ km ²)	low/mid	14	+0.913	0.036

Footprint sharing, not basin size, is the stratifier that tracks the linear neighbor effect. Among basins GRACE comfortably resolves, the high-sharing stratum carries +1.16% – the same magnitude as in the small basins – while the low-sharing majority runs *negative*. The statement is not absolute: one cell resists the pattern, the 14 small, nominally low-sharing basins, which also come in positive (+0.913%, $p = 0.036$). It is the smallest cell in the table, and it is the one our contamination metric is least able to score, since the metric counts only foreign *land* and therefore understates footprint sharing for coastal and island basins (Sect. 3.4) – which is what most small, nominally clean basins are. We flag it rather than explain it away. Repeating the cross at CSR’s published 200,000 km² caution threshold (Sect. 2) sharpens rather than weakens the picture: basins above the cut with high sharing keep the effect (+1.10%, $p = 0.033$), the low/mid-sharing majority above the cut is significantly *negative* (−0.80%, $p = 0.014$), and the resisting small-and-clean cell shrinks to marginal (+0.69%, $p = 0.057$). Contamination terciles reach the same conclusion without any threshold choice (low −0.27%, mid −0.18%, high +1.16%, $p = 0.00051$). The pocket is short-horizon as well as sharing-confined: for the ERA5-conditioned neighbor ridge the high-sharing tercile gives +1.11%




Figure 3. The lead-1 neighbor estimate against its null models. (a) Pooled RMSE of the real correlation graph (line) against the distributions of 50 degree-matched random graphs and the IAAFT surrogate pipelines. (b) Skill versus minimum allowed neighbor distance (no restriction, ≥ 300 km, ≥ 500 km, ≥ 1000 km): unchanged at 300 km (ruling out the crudest form of mascon leakage), halved at 500 km, and negative at 1000 km – bounding the spatial scale of whatever is being picked up. (c) The estimate without conditioning and with full 11-variable ERA5 conditioning: unmoved (+0.308 versus +0.322%), and non-significant in both cases – the confidence intervals cross zero, which is the point of the section.

at lead 1 ($p = 2.3 \times 10^{-4}$) and +0.59% at lead 2 ($p = 0.037$), then nothing at leads 3–6, while the mid and low terciles are negative at every lead.

Short-horizon, confined to shared footprints, and absent where footprints are clean is the signature of leakage, not of hydrology. A centroid-distance exclusion does not remove shared mascons for such basins, so we rest no novelty claim on the linear pocket, and we retire the “small basins” framing that stood in for it. Accordingly we describe this section as a characterization with null-model controls, not as isolation and validation of a signal: the controls tell us what the pooled number is *not* (chance graph structure, shared weather), the paired test tells us it is not distinguishable from zero, and the stratification tells us that the little that is there sits exactly where an instrument artifact would put it. Whether the same is true of the nonlinear correction is a separate question with a different answer, and Sect. 4.5 tests it directly.

4.4.3 Conditioning does not explain the effect away.

Giving every model in the comparison the 11 ERA5 local-meteorology variables barely moves the lead-1 point estimate (+0.322% conditioned versus +0.308% unconditioned), so whatever the estimate reflects, it is not a proxy for shared local weather. Both estimates are non-significant, however ($p = 0.115$ and $p = 0.199$), so the honest statement is that conditioning does not explain the effect away – not that the effect survives conditioning.

The controls constrain what the lead-1 estimate could be without ever lifting it clear of zero. It is not chance graph structure – the real graph beats all 50 degree-matched placebos – and it is not shared local weather, since full ERA5 conditioning leaves

Table 4. The lead-1 neighbor estimate against its null-model battery. Each row states the control, what it tests, and the outcome. All effects are pooled skill versus the own-basin reference on identical rows. The estimate being controlled is itself not significantly different from zero (+0.31 %, $p = 0.199$), so these rows bound what the estimate could be rather than validating a demonstrated effect.

Control	Question	Outcome
Degree-matched random graphs (50)	Chance / capacity?	passed at lead 1 only: beats 50/50 draws ($p_{\text{rank}} = .02$); at leads 2–6 beaten by 0/50
IAAFT surrogates (99)	Univariate spectra rather than cross-basin timing?	passed at lead 1: beats 99/99 surrogates ($p_{\text{rank}} = .01$); at leads 2–6 the real arm trails the surrogate distribution, matching the null there
≥ 300 km source exclusion	GRACE mascon leakage?	unchanged (+0.32 %, 50/50 placebos)
$\geq 500 / \geq 1000$ km source exclusion	How far does it reach?	decays then dies (+0.16 % at 500 km; -0.60 % at 1000 km, beaten by 0/50 placebos)
ENSO + IOD conditioning	Shared teleconnections?	point estimate unmoved (+0.310 % vs. +0.308 %); still ahead of all 50 draws; the pooled contrast is non-significant either way
Full ERA5 conditioning (11 variables)	Shared local weather?	point estimate unmoved (+0.322 % vs. +0.308 %); both non-significant
Train-only 6σ jump screen	A few outlier basins?	2 flagged; screened estimate +0.21 %, still 50/50 vs placebos
Area stratification ($\geq / <$ 90,000 km ²)	Small-basin leakage receivers?	not passed: the whole pooled effect sits in the 35 small basins (+1.03 %, $p = .0007$) and none in the 199 large ones (+0.10 %, $p = .72$)
Footprint-sharing 2×2 (area $\times$ contamination)	Is it size, or is it shared mascon footprint?	not passed, and it is sharing not size: high-sharing basins carry +1.16 % whether or not GRACE resolves them; low-sharing resolved basins run -0.50 %

All rows are from the corrected issue-date pipeline; archived source files are listed in Appendix A.

it unmoved. The centroid-distance exclusions bound its spatial scale ($\sim 300\text{--}1000$ km). But the footprint-sharing stratification places what is left exactly where an instrumental explanation predicts it, and a centroid-distance rule between basin *centroids* does nothing about two basins drawing on the same mascon. The reading most consistent with the whole battery is that at the linear tier the neighbor is largely acting as a second view of partly the same measurement – which would explain why the estimate is confined to shared-footprint basins and to leads 1–2, and why it is negative beyond 1000 km. We claim no hydrological transfer, and at the linear tier we claim no effect.




Figure 4. Geography of the lead-1 neighbor estimate. Global choropleth of the per-basin Diebold–Mariano statistic for the neighbor ridge versus its own-feature twin (234 basins; blue: neighbor helps, red: neighbor hurts), with basins passing Benjamini–Hochberg FDR control ($q = 0.10$) hatched following Wilks (2016). The map is two-sided by construction and should be read alongside the pooled null result of Sect. 4.4: a stable geography of helped and hurt basins is compatible with no pooled effect. Twenty-one of the 234 basins survive FDR control: the neighbor helps significantly in 13 and hurts significantly in 8. The extremes are as divided as the pooled null implies. Among the helped are the East Brazil South Atlantic coast ($DM = -3.80$), the endorheic Tarim He–Lop Nur basin (-3.57), and the South Hudson Strait–Ungava Bay coast (-3.38); among the hurt are the Yenisey ($+4.26$), the South Gulf of Oman–West Arabian Sea coast ($+4.00$), and Madagascar ($+3.78$).

4.4.4 Heterogeneity, and per-basin selection.

A pooled null does not mean a spatially uniform null: a stable pattern of helped and hurt basins can average to nothing. Two analyses probe this – a per-basin FDR map (Fig. 4) and a leave-one-fold-out selection experiment that applies the neighbor only where it helped in the other four folds. The latter is a measure of cross-fold stability in per-basin benefit, not an operational recipe: selection for early folds uses folds lying later in time, so a deployable version would have to select on past folds only. On the corrected pipeline the per-basin decomposition confirms the two-sided structure: 128 of 234 basins nominally favor the neighbor and 21 survive FDR control at $q = 0.10$ – **13 helped and 8 hurt**. The hurt set is not noise: it is the familiar great-river roster (Yenisey, Paraná, Yangtze, Amur) whose large, clean footprints invite spurious distant neighbors. The helped set skews African and coastal (Niger, São Francisco, Rukwa, Namibia, Zambezia, the West Nile Delta). Static covariates barely predict who benefits (random-forest cross-validated $R^2 = 0.05$; the top predictors are basin size and cross-fold neighbor stability), but past per-basin performance does: applying the neighbor only where it helped in the other four folds lifts the pooled lead-1 effect from $+0.31\%$ (n.s.) to $+1.12\%$ ($p < 10^{-9}$, ~ 128 basins selected), 63 % of the in-sample oracle’s $+1.78\%$, with the same pattern on the ERA5 backbone ($+1.05\%$). Under the caveat above, this is a stability measurement, not a deployable recipe. We do not restate the pre-correction values, because the pooled effect they decomposed no longer exists.

4.4.5 Architectures that should have enlarged the effect do not.

Predictive-lag neighbor selection loses to the own-basin reference at every lead. Alternative graph constructions and structurally richer filters – predictive-lag selection, a single-latent fusion filter treating the neighbor as a second sensor of the same state, and the bivariate filter with correlated process noise – were all built and all failed to enlarge the effect. On the corrected pipeline, predictive-lag selection is -1.30% against the own-basin reference at lead 1 ($p = 0.0007$) and beats 0 of 50 placebo graphs – worse than a random basin; the fusion filter is -2.14% against the plain Kalman ($p = 6 \times 10^{-5}$; its MLE-weighted variant -6.27%), likewise 0/50; and the coupled bivariate filter is indistinguishable from the plain Kalman (-0.02% , $p = 0.96$), ahead of 45 of 50 placebo draws but short of significance ($p_{\text{rank}} = 0.12$). Two-hop neighborhoods are no better than one-hop on the corrected pipeline (residual MLP $+0.13/-0.28/-0.15\%$, all n.s.; the GAT is significantly worse at lead 2), with one disclosed exception: the ERA5-conditioned residual MLP prefers two hops at lead 2 ($+0.82\%$ over one-hop, $p = 0.006$, all seeds positive), though chained against the own-basin twin the residual two-hop gain is only $\sim +0.2\%$. Together with the GAT null (Sect. 4.2), the linear-tier conclusion is that no standard linear or message-passing construction extracts usable cross-basin information from observed TWSA fields. That conclusion holds until the delivery mechanism changes, which is the subject of the next two subsections.

4.4.6 One nonlinear exception, and it is neighbor-specific.

The linear null is not the whole story even at the GRACE-only information set. A two-stage residual MLP fed *only* the neighbor’s propagated filtered state beats its ridge twin on identical features by $+0.81$, $+2.15$, and $+2.17\%$ at leads 1–3 (three-seed ensemble, $p \leq 6.7 \times 10^{-5}$, every seed individually significant). The control that matters is the placebo family: every seed beats all 20 degree-matched random graphs at every horizon (9 of 9 seed–horizon cells: three seeds, three leads), and the placebo family itself pools to within 5×10^{-4} RMSE of the plain own-basin ridge. A junk neighbor through the same two-stage architecture adds nothing; the real neighbor adds up to two percent. This is the one place where cross-basin information survives on the same GRACE-only feature set that produced the linear null. The signal is nonlinear only, and it disappears – turning significantly negative – as soon as ERA5 forcing enters the model (-0.67 , -0.62 , -0.55% , all $p < 0.05$). It supports the reading that cross-basin information matters most where forcing data are absent.

4.5 Representation decides: the neighbor is usable only as a propagated state

This experiment holds the information *source* fixed – the same selected neighbor’s Kalman-filtered state – and varies how it reaches the model. Two things vary in the headline pair, and separating them turns out to matter more than we expected. The input-channel arm receives the neighbor’s 12-month filtered-state history as a full encoder channel, while the correction stage receives a single number per forecast, the neighbor’s ρ^h -propagated current state. Delivery and representation are therefore confounded in that pair, so we ran the arm that separates them: the correction stage fed the identical 12-month history. The answer is that representation, not delivery, is the operative variable, and the rest of this section is organized around that. Feeding the neighbor’s filtered-state history to the LSTM as one more encoder input channel yields no robust ensemble gain (Fig. 5):

Table 5. The stage-2 neighbor correction: skill (%) of the stacked system versus its own stage-1 LSTM (same network within seed; the contrast isolates the correction). The ensemble column is the reported quantity – predictions averaged across the two seeds per row, then scored – and the per-seed columns are the robustness check. Confidence intervals are moving-block bootstrap on the ensemble. The last column gives the number of 20 degree-matched random-graph correction placebos beaten, per seed.

Lead	Seed 0	Seed 1	Ensemble	95 % CI (ens)	DM p (ens)	Placebos beaten
1	+0.64	+1.19	+0.91	[+0.64, +1.22]	1.1×10^{-10}	20/20, 20/20
2	+1.33	+1.56	+1.45	[+1.07, +1.91]	2.3×10^{-10}	20/20, 20/20
3	+1.29	+1.36	+1.33	[+1.02, +1.69]	1.9×10^{-9}	20/20, 20/20
4	+1.26	+1.37	+1.32	[+0.96, +1.76]	3.1×10^{-8}	20/20, 20/20
5	+1.51	+1.28	+1.42	[+1.06, +1.84]	6.2×10^{-9}	20/20, 20/20
6	+1.90	+1.93	+1.96	[+1.66, +2.36]	1.2×10^{-13}	20/20, 20/20

Placebo graphs are redrawn independently for each fold and horizon, so the twelve cells of the last column are separate draws; with 20 draws the rank statistic is nonetheless pinned at its $1/21$ floor (Sect. 5.5). Archived source files are listed in Appendix A.

at lead 1 the two-seed ensemble gains +0.36% ($p = 0.090$), and at leads 2–6 it is negative throughout (−0.42, −0.43, −0.61, −0.37, −0.38%; $p = 0.038, 0.067, 0.010, 0.169, 0.332$). We state this as “no robust ensemble gain as an input channel”, and no more than that. The channel is not shown to be useless, and it is not shown to be damaging: the per-seed picture is unstable enough that at lead 2 the two seeds disagree in sign, which is exactly why the ensemble is the reported quantity and why an earlier single-seed reading of this experiment was an artifact. The representation-matched arm settles what that failure is about. Giving the correction stage the input arm’s 12-month history in place of the propagated scalar – same stage, same capacity, same stage-1 network, only the representation changed – removes essentially the whole effect: the two-seed ensemble scores +0.05, +0.35, −0.42, −0.88, −0.78, and +0.25% at leads 1–6, not significant at any lead, and its placebo ranks collapse (2 of 20 beaten at lead 3 in one seed). Against the propagated-scalar correction directly it loses at every lead (−0.87 to −2.24%, all $p \leq 5.0 \times 10^{-3}$, both seeds). That is the same place the input-channel arm sits. Handed the neighbor’s raw history, neither architecture extracts anything; handed the ρ^h -propagated state, the correction stage extracts +0.9 to +2.0%. We therefore do not claim that the correction stage is what finds the signal. What the neighbor contributes is usable only once it has been propagated to the forecast target time, and the propagation is supplied by the filter of Sect. 4.1, not learned here. The honest statement of the delivery experiment is that a deep encoder given 12 months of a neighbor’s history does not recover, at this sample size, the alignment that one line of AR(1) propagation provides for free.

The stacked system trains the identical LSTM (own state + ERA5, no neighbor) and then fits a small MLP to the stage-1 training residual from the propagated neighbor state alone. Because challenger and reference share the same stage-1 network within each seed, the contrast isolates the correction. Table 5 shows the result: the same information delivered this way adds +0.91 to +1.96% at every lead 1–6, in both seeds, with every confidence interval clear of zero.

Six properties give the correction its evidential weight.




Figure 5. Representation decides whether the neighbor signal survives a deep model. Skill increment (%) of the neighbor information over the same own-state + ERA5 LSTM, by lead, for the two routes: as an additional encoder input channel (two-seed ensemble +0.36% at lead 1, negative at leads 2–6, no lead significant at the 1% level) and as a dedicated residual-correction stage (+0.91 to +1.96% at every lead 1–6, both seeds). Shaded band: the range of 20 degree-matched random-graph placebo corrections riding the matching stage-1 network, whose mean cost is -0.27 to -0.05% – the pure variance price of adding a second stage. The real correction sits 6.8 to 18.6 placebo standard deviations beyond that band.

First, and primarily, **the random-graph placebo**. This is the control the claim rests on. Its logic is that a placebo neighbor is a feature the stage-1 network has never seen – novel, like the real neighbor – but with the wrong identity, so it isolates identity from novelty. The correct statistic is the *increment* each stacked model adds over its own stage-1 network, compared with 20 random-neighbor corrections riding the matching stage-1 network. Placebo increments cost only -0.27 to -0.05% – a junk-neighbor correction buys nothing and pays the variance price of a second stage, which the audit’s error decomposition predicts to within a few hundredths of a percent – while the real increment beats all 20 draws in all twelve lead–seed cells. We report the margin rather than the rank p , which is pinned at its $1/21$ floor: the real correction sits 6.8 to 18.6 placebo standard deviations beyond the placebo mean. Measured against the placebo median rather than against stage 1, the increment is $+0.92$ to $+1.96\%$ – two independently constructed references agreeing to within 0.1–0.3 percentage points at every lead. An earlier version of this test showed a single blemish (16 of 20 at lead 2 for seed 1); it was an artifact of scoring a seed-1 model against placebos riding the seed-0 network, and it disappeared once each seed was scored against its own placebos. One accounting caveat, which we do not minimise: the draws are redrawn per fold and horizon, but the two seeds within a lead are scored

against the same draws, so the twelve lead–seed cells are six independent draw sets each scored twice, not twelve independent events.

Second, **the own-state control, a diagnostic rather than the primary control.** This neighbor-free variant feeds each basin’s *own* propagated state through the identical feature path, same stage-1 network, same stage-2 model, capacity, and seed. Mechanically it is exact: the self-graph reproduces each basin’s own propagated state to $\max|\Delta| = 0.000 \times 10^0$, and no zero-fill asymmetry exists between the two variants, since all 234 basins have a top-1 neighbor in every fold. The control actively harms: -1.16 to -2.12% across leads ($p \leq 2.5 \times 10^{-7}$). We use it as evidence that the stage-2 head has ample capacity to do damage, so the stacked system’s gain is not a capacity artifact, and as evidence for the in-sample residual caveat below. We do *not* use the neighbor-versus-own-state contrast as a skill number or as an “architecture term”, because the control is too easy: its correction is anti-correlated with the stage-1 error (-0.08 to -0.16), and only 0.16 – 0.54 percentage points of its harm is variance cost, meaning 79 – 85% of it is a systematically wrong-direction correction learned from in-sample residuals. A control that is actively harmful inflates any contrast drawn against it.

Third, **agreement between the two defensible references.** The correction is worth $+0.91$ to $+1.96\%$ against its own stage-1 network and $+0.92$ to $+1.96\%$ against the placebo median. Independent constructions landing on the same interval is the strongest single property of this result. Fourth, **the correction is positive in every fold as well as every stratum:** decomposing the ensemble increment over the five evaluation folds gives a positive value in all 30 fold–lead cells (minimum $+0.25\%$, in the earliest-freeze fold at lead 1; maximum $+2.63\%$), so no single evaluation period carries the result. At the basin level, 22 of 234 basins pass FDR control at lead 1 – 20 helped, 2 hurt, against 13 and 8 for the linear estimate – and the stack’s per-basin DM map rank-correlates with the linear neighbor map at $\rho = +0.49$ ($p = 1.4 \times 10^{-15}$): the two tiers largely agree on *where* the neighbor helps, and differ in how much of it they convert into skill.

Fifth, **horizon stability that no other component has.** The stage-1 gain over the lag-feature ridge is gone by leads 5–6, where the LSTM is indistinguishable from that ridge (-0.04% , -0.11%). The correction does the opposite: it holds at $+1.42$ and $+1.96\%$ there, with the lead-6 value the largest in the column. At long leads it is the only architectural gain in the study that has not decayed.

Sixth, **it does not carry the leakage signature.** The obvious objection to any cross-basin claim on GRACE is that the “neighbor” is partly the same measurement, and Sect. 4.4 showed that this is exactly what the linear pocket looks like. We therefore put the stacked correction through the identical stratification – contamination terciles, the $90,000 \text{ km}^2$ area split, and their 2×2 cross – at all six leads. The correction is positive in *all 60* strata-by-lead cells, and it does not concentrate in the strata a leakage artifact would favour:

Stratum	Basins	$h=1$	$h=2$	$h=3$	$h=4$	$h=5$	$h=6$
Sharing, low tercile	78	+0.69	+1.31	+1.29	+1.28	+1.31	+1.84
Sharing, mid tercile	78	+1.09	+1.58	+1.44	+1.33	+1.61	+2.01
Sharing, high tercile	78	+1.00	+1.48	+1.29	+1.35	+1.37	+2.03
Resolved $\times$ low/mid sharing	142	+0.89	+1.57	+1.46	+1.38	+1.49	+2.00
Small ($<90,000 \text{ km}^2$)	35	+0.64	+0.88	+0.93	+0.82	+0.77	+1.05

All cells in the first four rows are significant ($p \leq 6.9 \times 10^{-3}$ in the low tercile, $p \leq 1.3 \times 10^{-4}$ elsewhere); the small-basin row is significant at every lead ($p \leq 0.044$) despite its n . The decisive cell is the fourth row: 142 basins that GRACE resolves and whose footprints are comparatively clean, where the *linear* effect is -0.50% and the correction is $+0.89$ to $+2.00\%$. The per-seed worst cell across every stratum and lead remains positive ($+0.22\%$, in a 21-basin cell).

625 “Positive everywhere” is not the same as “the same everywhere”, and along the size axis it is not. Small basins gain *least*, and the gradient is real rather than visual: the size interaction is significant at every lead (lead-1 $p = 0.026$, lead-6 $p < 0.001$). We state that plainly because its direction is the reverse of the artifact’s prediction – leakage would make the smallest, most footprint-shared basins gain *most*, and they gain least. Along the contamination axis the differences are not statistically distinguishable at any lead (high-vs-low interaction $p = 0.099$ at lead 1, with the point estimate trending toward *more* gain in
630 high-sharing basins, and $p \geq 0.51$ at every other lead), which turns “not concentrated” from a description of the table into a tested statement.

The linear pocket and the stacked correction therefore have different signatures: one is confined to shared footprints and to leads 1–2, the other is positive in every stratum at every lead and, if anything, weakest where the artifact would be strongest. **Under this stratification the correction shows no detectable concentration where leakage would put it, and the size
635 gradient runs opposite to the artifact’s prediction.** That is the strongest statement the design supports: the contamination metric is a lower-bound proxy (it counts only foreign land), the interaction test is a null result rather than proof of absence – with the lead-1 point estimate trending toward more gain under high sharing ($p = 0.099$) – and the stratification reuses the headline’s test months. Ruling on footprint sharing aside, the analysis does not identify what the neighbor is contributing, and we make no claim that it is a hydroclimatic teleconnection.

640 A diagnostic points the same way. The ERA5 residual MLP with neighbor features concatenated into its inputs – rather than given their own stage – is negative overall (-0.67 , -0.62 , -0.55%), and its harm is concentrated in the *low*- and mid-sharing basins (-0.89% , $p = 0.011$, in the resolved low/mid stratum) while being indistinguishable from zero where footprints overlap. Naive concatenation imports noise except where the neighbor is partly the same measurement; only the dedicated stage turns neighbor state into gains that are positive across every stratum.

645 Why does the same neighbor signal succeed in one form and not the other? Our first reading was gradient starvation (Pezeshki et al., 2021): in joint training a strong feature set (the forcing sequence) starves the gradient signal of a weak, partially correlated feature (the neighbor state), so the weak channel is ignored or fit to noise, and a dedicated residual stage rescues it by giving the weak feature its own objective. The representation-matched arm rules that out as the whole story. A dedicated stage fed




Figure 6. The linear neighbor effect carries a leakage signature; the stacked correction does not. (a) Linear tier: skill of the ERA5-conditioned neighbor ridge over its own-state twin by mascon-contamination tercile ($n = 78$ each), leads 1–6 – significant only in the high-sharing tercile and only at leads 1–2; everywhere else the point estimates are negative or indistinguishable from zero. (b) Stacked tier: the stage-2 neighbor correction over its own stage-1 LSTM (two-seed ensemble) for the same terciles plus the resolved, low/mid-sharing stratum ($n = 142$), with moving-block bootstrap ribbons – positive in every stratum at every lead, including where the linear effect is negative. Both panels share the y-axis; the stratification reuses the headline’s test months (Sect. 5.5).

the neighbor’s raw 12-month history has its own objective and nothing else to fit, and it still recovers nothing (Sect. 4.5); if
650 training dynamics were the binding constraint, that arm should have worked. Gradient starvation may still contribute to the
input channel’s instability, but it cannot explain a failure that survives the separation of objectives.

What distinguishes the arms that work is that the neighbor arrives already propagated to the forecast target time. $\rho_j^h x_j(t)$ is
not a compression of the history; it is the neighbor’s state carried forward by its own estimated dynamics, so the downstream
model is handed a quantity that is temporally aligned with what it must predict. Recovering that alignment from 12 raw monthly
655 values is a learning problem in its own right, and at this sample size – a few thousand rows per fold, one weak predictor among
many – it is evidently not solved. The reading we can support is that the filter does more than clean the target: it also puts cross-
basin information into the only form in which our models can use it. That places the propagation, not the architecture, at the
center of the third contribution, and it is a claim about sample size and inductive bias rather than about optimization pathology.
The residual-MLP experiments of Sect. 4.4 are consistent with it: they also feed propagated states, never raw history, where a
660 neighbor-only residual MLP over the plain ridge beats its ridge twin on identical features by +0.81 to +2.17% while beating
all 20 placebos in every seed–horizon cell – and it does so despite the *linear* neighbor contrast on that same information set
being a null. Two architectures, two information sets, one conclusion: the cross-basin signal is there, and it is reachable only
when the neighbor is presented as a propagated state. A double-counting check completes the picture: stacking the correction
on top of an encoder that *already ingests* the neighbor history yields -0.46% at lead 1 before drifting weakly positive at long

665 leads (+0.23 to +0.83 %), never approaching the +0.91 to +1.96 % the same correction earns over a neighbor-free encoder – consistent with the representation account: once the encoder has consumed the raw history, a second helping of the propagated state has little left to add, and the clean two-stage design remains the better construction.

4.5.1 Required disclosure on model ranking and selection.

The stacked system is the nominal best model at leads 4–6 among all arms evaluated there, and against the lag-feature linear
670 system with the same predictor *types* it gains +2.75 to +3.86 % at every lead. It is *not* the best model at short leads once history length is equalized: the flat 12-month ridge of Sect. 4.2 ties it at lead 1 (−0.29 %, $p = 0.58$) and beats it at leads 2–3 (−1.31 %, $p = 8.3 \times 10^{-4}$; −0.77 %, $p = 0.032$), and that arm exists only at leads 1–3, so the long-lead ranking is untested against it. A leaderboard position is not the evidence, however, and one comparison in particular must be kept in view: the largest architectural number anywhere in this study is *neighbor-free* – a two-stage residual MLP on own-basin ERA5 features
675 beats its ridge twin by +3.35, +3.58, and +2.84 % at leads 1–3. Two-stage nonlinear correction is therefore worth more, in absolute terms, than the neighbor correction is, and a cross-architecture ranking cannot separate the two. On the corrected pipeline the direct contrast confirms it: the stacked system and the neighbor-free resMLP-ERA5 ensemble are a statistical tie at all three shared leads (−0.30, +0.19, −0.10 %; DM $p = 0.49, 0.63, 0.81$). The evidence that the neighbor adds skill is the controlled contrast of Table 5 – same stage-1 network, placebo-verified increment – not the cross-architecture leaderboard,
680 and any deployment claim should be read accordingly. Moreover, the stacked architecture was chosen after inspecting earlier architectures’ results on these same five folds, and no untouched evaluation period exists within the GRACE-FO record; its results are therefore **exploratory**, and the reported p values do not account for architecture-selection uncertainty. The stage-2 increment is protected by its seed-matched placebo battery, but confirmatory status awaits an untouched period or an independent product.

685 5 Discussion

5.1 What the benchmark result means, and does not mean

The benchmark result is an argument about reference classes, not a new forecasting method. Persistence-type references are supposed to represent “no information beyond memory”; for a noisy observable they leave the persistence class’s own headroom unexploited, and that headroom on deseasonalized GRACE basin series is 5.0–8.8 % at the leads where most data-driven
690 papers claim their skill. Any study that benchmarks against persistence, damped persistence, or climatology on observed TWSA – i.e., most of the literature cited in Sect. 1 – overstates the information its model has extracted by roughly this margin, whatever mixture of measurement error and unmodeled dynamics the filter is exploiting (Sect. 3.3). The fix costs three parameters per basin. This is the same benchmark-hygiene move that Niraula and Goessling (2021) made for the sea-ice edge via spatial damping, executed here via signal/noise separation, and it converges with three independent lines: persistence-class methods
695 underrated at the coefficient level (Kankanige et al., 2026), statistical per-basin prediction competitive at basin level (Ahi and

Cekim, 2021), and linear models unbeaten when autoregressive information is excluded (Nie et al., 2025). Our contribution to that conversation is an ablation-grounded attribution: with autoregressive information included, what separates the models the field builds is mostly how well they filter it – deleting only the noise term from the benchmark deletes the margin (Sect. 3.3). The ablation fixes the attribution within the model class; what the filtered component physically is – instrument noise, fast hydrology, decomposition error – remains unidentified (Sect. 3.3).

Two scope notes. The argument applies to *observed* TWSA; targets that are themselves model output, reconstructions, or gap fills (e.g., the HydroGlobe target of Nie et al., 2025, or the products of Humphrey and Gudmundsson, 2019; Yin et al., 2023; Mo et al., 2022) contain no satellite observation noise to filter, so neither the benchmark inflation nor its fix transfers. And our filter is not data assimilation (Springer et al., 2026): nothing is assimilated into any model, and no state but the basin’s own scalar signal is estimated. The closest upstream machinery we know of is the integrated postprocessing of Zhang et al. (2025), which models the non-seasonal GRACE signal as a stochastic process and iteratively separates it from noise in the monthly solutions. That work shows signal/noise separation improves *extraction* of TWSA from the existing record; ours differs in pointing the separated state forward – using it to forecast, and to benchmark forecasts – so the benchmark claim here is not that filtering GRACE is new, but that forecast comparisons on observed TWSA are uninterpretable without it.

5.2 The crossover as a design map

The initial-conditions-versus-forcing partition is old (Wood and Lettenmaier, 2008; Shukla and Lettenmaier, 2011; Koster et al., 2010; Yossef et al., 2013; Pechlivanidis et al., 2020); what is new is its location on satellite-observed TWSA, measured between two published data-driven systems with paired tests. The crossover at ~ 2 months is earlier than for streamflow in snow-dominated regions but consistent with TWSA’s strong memory being concentrated in exactly the component the filter extracts. The practical reading is a design map, not a contest: below the crossover, invest in filtering the initial condition (three parameters suffice; further architecture buys little); above it, invest in exogenous forcing and in a longer window of the forcing history – our own evidence (Sect. 4.2) shows a ridge on issue-time forcing lags cannot carry ERA5 skill past lead 1, models reading a 12-month forcing window carry it to lead 4 (our sequence model’s gain over the lag ridge is significant through lead 4 and gone by lead 5, though whether that last step needs the sequence architecture rather than the longer window is untested – the flattened-history ridge exists only at leads 1–3), and Li and Kusche (2026)’s system, which commits fully to this side, carries forcing-driven skill to lead 6 and beyond. That the linear lag-feature ERA5 gain stops at lead 1 sharpens rather than weakens the map: the useful part of forcing information at longer leads is not in the issue-time lags but in the depth of history a model reads. At basin level the two edges are uncorrelated (Spearman -0.141 between GRACE-FCast’s lead-4 edge and the neighbor benefit), so the complementarity holds across basins as well as across leads. A genuinely hybrid product – filtered initial conditions, forcing-driven long leads, and a spatial correction – is the obvious system nobody has built; our splice is a descriptive sketch of its headroom, not that product.

5.3 The neighbor signal in context

Spatial context is increasingly assumed to help TWSA prediction (Li and Kusche, 2026; Arzoumanidis et al., 2026; Steidl and Zhu, 2025). Our results discipline that assumption more than they support it, and the two-sided answer is the contribution. On the discipline side, the linear result is a *null*: one correlation-selected neighbor adds +0.31 % at lead 1 ($p = 0.199$) and nothing at longer leads. Worse for the assumption, what little sits inside that null has a leakage signature – it is confined to basins whose mascon footprint is shared with land outside the basin and to leads 1–2, and it is *negative* in the 142 resolved, clean-footprint basins (Sect. 4.4). A graph attention network given identical information never beats a ridge, and two-hop neighborhoods never beat one-hop. Anyone building a basin graph for observed TWSA should expect these to be the default outcomes, and should stratify by footprint sharing before believing a positive one.

On the support side, the signal is not absent – it is inaccessible to the standard tools. Two nonlinear, two-stage constructions recover it: a neighbor-only residual MLP over a plain ridge at the GRACE-only information set (+0.81 to +2.17%) and a neighbor-only correction stage over a deep own-state + ERA5 LSTM (+0.91 to +1.96% at every lead 1–6). Both beat their full random-graph placebo families in every cell, so what is recovered is the neighbor’s identity and not the extra stage. And the correction does not inherit the linear pocket’s signature: it is positive in all 60 strata-by-lead cells, shows no statistically distinguishable difference across footprint-sharing terciles, and is worth +0.89 to +2.00 % in precisely the resolved, clean-footprint stratum where the linear effect is negative. Along the size axis it does vary, in the direction that helps rather than hurts the reading: the smallest basins – the ones leakage would favour most – gain the least. That rules out the leading artifact explanation. It does not identify a mechanism, and we advance none: what the two-stage delivery extracts from a neighbor basin’s history remains an open question, and calling it a teleconnection would be an overclaim. Steidl and Zhu (2025) forecast with a 10-nearest basin graph but on a reconstructed target, with no persistence-class baselines, no isolation of the graph’s contribution, and no null models; our GAT null, control battery, and footprint-sharing stratification are the coherent companion measurement – the graph assumption, tested, yields a linear null with a leakage pocket, a small nonlinear signal that survives the leakage stratifier, and the finding that message passing is the wrong way to extract it.

A regional reading of where the neighbor channel pays would be valuable in a water-security context (Fig. 7), particularly the hypothesis that spatial TWSA information matters most where reanalysis forcing is least trustworthy. The corrected pipeline supports it from three directions. Per basin, ERA5 gain and conditioned neighbor gain are anti-correlated (Spearman -0.220 , $p = 0.0007$, 234 basins). By continent, the two are mirror images: Africa has the smallest median ERA5 gain (+0.7%) and the largest median neighbor gain (+3.1 % conditioned; median per-basin skill +1.6 % on the linear tier), while Europe shows the reverse (+7.6 % ERA5, -1.9 % neighbor). And pooled Africa sharpens rather than dilutes under conditioning: the unconditioned African neighbor effect is +0.85 % ($p = 0.15$), unchanged by index conditioning, but rises to +1.57 % ($p = 0.0035$) once all eleven ERA5 variables are aboard – spatial TWSA information is worth most precisely where the reanalysis has least to say. The GRACE-only residual-MLP result points the same way (Sect. 4.4).




Figure 7. Do the two information sources cover each other’s dead zones? Scatter of per-basin lead-1 ERA5 gain against per-basin lead-1 ERA5-conditioned neighbor gain (234 basins, colored by continent), with continental pooled values inset. The hypothesis the panel tests is that ERA5 gains concentrate in the extratropics and fail in the humid tropics, while any neighbor benefit concentrates where reanalysis forcing is weakest. On the corrected pipeline the rank correlation is -0.220 ($p = 0.0007$); continental medians run from Africa ($+0.7\%$ ERA5, $+3.1\%$ neighbor) to Europe ($+7.6\%$ ERA5, -1.9% neighbor).

5.4 Comparability of published numbers

No published study reports pooled deseasonalized basin RMSE over a global sample, so no external number is directly comparable to ours; we compare only against baselines recomputed under our protocol, and we flag two specific numbers a reader might juxtapose. Mo et al. (2025) report median grid-cell RMSE of 1.79/2.07/2.26 cm at 1–3 months; that figure is a median over grid cells of *full-signal* errors (seasonal cycle included) on a different window, three choices that each shrink it relative to our pooled deseasonalized basin metric, and only their code – not their predicted fields – is archived, so no matched-protocol comparison is possible. Li et al. (2026) report ~ 1 cm domain-averaged deseasonalized RMSE for Africa; spatial averaging over the domain removes most variance, and their basin-level skill is categorical (ROC), not RMSE. We also note the FLDAS-Forecast hindcast fields are public (<https://ldas.gsfc.nasa.gov/fldas/models/forecast>); we did not add a third head-to-head because the product is regional (its Africa/Middle East domain covers ~ 37 of our 234 basins), percentile/ROC-oriented, and ends in December 2020, overlapping only ~ 19 months of our test targets. Its published conclusion – that initial-condition persistence dominates sub-seasonal TWS skill (Li et al., 2026) – is the qualitative claim our benchmark formalizes.

5.5 Limitations

1. **Stage-2 residuals are in-sample, and this is now measured rather than assumed.** The headline stacked correction is trained on the LSTM’s training-set residuals. The *evaluation* was always honest (shared stage-1 network, no test information reaches the MLP), but because the LSTM overfits its training set we could only argue that the learned
775 correction was plausibly attenuated. We have since rerun the correction on out-of-fold residuals, fitting stage 2 on three contiguous blocks of issue months held out of the stage-1 fit. The out-of-fold correction is *larger* at every lead (+1.14 to +2.28%, against +0.91 to +1.96% in-sample; the direct contrast is +0.20 to +0.34%, significant at all six leads) and beats 20 of 20 placebos in every cell. The in-sample protocol was therefore conservative, and the headline numbers we report understate the effect. The residual caveat that remains is the one this mechanism explains rather than inflates: it is
780 why the own-state control learns a systematically wrong-direction correction (Sect. 4.5).
2. **The rank statistic is pinned at its floor.** With 20 placebo draws per cell the rank p cannot fall below $1/21$, which is why we report placebo standard deviations alongside it. Draws are now redrawn independently per fold and horizon, so the twelve lead–seed cells of Table 5 are no longer correlated re-scorings of one set of 20 realizations, as they were before the 2026-08-15 repair.
- 785 3. **Controls not yet ported to the stacked system.** The ≥ 300 km exclusion and meteorology conditioning that bound the leakage and shared-climate confounds were run for the linear estimate, not for the stacked correction; the stacked evidence rests on the seed-matched placebo graphs alone. Indirect, and disclosed as such.
4. **Comparison asymmetries with GRACE-FCast.** Their training uses the Yin et al. (2023) reconstruction across the mission gap (hindsight); their seasonal/trend handling contributes ~ 0.07 – 0.09 of the full-product edge at leads 4–6
790 (bounded by the non-seasonal variant, which still wins those leads); our lead-1 win partly reflects the Kalman filter being tuned to exactly this target definition; and the matched sample excludes fold 5 and seven island basins.
5. **Splice numbers are post hoc** (Sect. 4.3): splice points chosen on the evaluated sample, several variants, no multiplicity control.
6. **Sample and era.** One mascon product (CSR); test windows all lie in the post-2019 climate era, and any leave-one-fold-
795 out selection, though honest in time, cannot escape that era; per-basin estimates rest on 83 lead-1 months.
7. **Seed budgets.** Two seeds for the LSTM stages, three for the residual MLP. Seed spread is comparable to the effects being measured – the LSTM sequence gain nearly halves from the better to the worse seed, the ERA5 residual MLP has a seed-2 outlier that flips sign at lead 1, and the LSTM neighbor increment at lead 2 flips sign between seeds. This is why every headline is a seed ensemble and no single-seed number is reported as a result.
- 800 8. **The filter’s noise parameter is not identified as measurement noise.** 259 of 1170 basin–fold fits push r to its lower boundary and 31 basins sit there in all five folds, and the AR(1)-plus-white-noise decomposition absorbs unmodeled dynamics into r (Sect. 3.3); an $r = 0$ ablation and replication on JPL/GSFC mascons are the concrete checks.

- 805 9. **Post-selection inference.** All architectures were developed and compared on the same five test folds; the stacked system combines winners of those comparisons, so its results are exploratory (Sect. 4.5) and nominal p values understate total uncertainty. No untouched GRACE-FO period exists for confirmation within this record.
10. **Shared mascon footprints carry what pooled linear neighbor effect there is.** The footprint-sharing stratification (Sect. 4.4) is the primary sensitivity for the linear neighbor result: among basins GRACE resolves, it places the pooled estimate in the shared-footprint stratum and turns it negative elsewhere, with one small low-sharing cell ($+0.91\%$, $p = 0.036$, 14 basins) flagged as the exception rather than explained away.
- 810 11. **The contamination metric is a lower bound, and the stratifications are not independent samples.** Contamination counts only foreign *land* – with land defined by CSR’s official v02 land mask and tile assignment by the official mapping file (Sect. 2.1) – so the ocean fraction of a coastal tile is not counted as foreign and island and coastal basins score low; true footprint sharing is therefore at least as large as we measure. The linear 2×2 is lead 1 on the neighbor ridge only; the stacked stratification covers leads 1–6 but reuses the same test months as the headline, so it is a decomposition of the
- 815 headline rather than an independent replication.

6 Conclusions

For basin-average, deseasonalized GRACE TWSA, the field’s standard persistence benchmark is weaker than its own reference class: a three-parameter per-basin state-space filter applied to the same observations beats damped persistence by 5.0–8.8 % at leads 1–2 and the per-basin ridge – the best conventional model in our ladder at short leads – at five of six leads, at negligible

820 cost. That margin is unchanged by two evaluation-protocol corrections that we applied to our own pipeline mid-study, which is the strongest evidence we can offer that it is a property of the benchmark rather than of our implementation. We propose it as the default reference for this task, and we read the consequence constructively: benchmarked properly, short-lead TWSA forecasting is revealed to be a filtering problem, and the meaningful question becomes what information beats state-space filtering.

825 The answer has three parts. Exogenous forcing beats it – ERA5 predictors add $+4.6\%$ at lead 1 from issue-time features alone, a twelve-month window of the same forcing – read linearly – adds $+3.3$ to $+5.0\%$ more at leads 1–3, beating every sequence model we trained on the identical history – and a published system built entirely on exogenous forcing (Li and Kusche, 2026) overtakes our best system beyond lead 2, locating an empirical crossing consistent with the classic initial-conditions-versus-forcing crossover on satellite-observed TWSA at about two months. Cross-basin information is the two-sided part. At

830 the linear tier it is a null – $+0.31\%$ at lead 1 ($p = 0.199$) and nothing at longer leads – and what little sits inside that null looks like an instrument artifact: stratifying on how much of a basin’s mascon signal comes from outside the basin places the estimate in shared-footprint basins at leads 1–2 across the size range, and turns it negative in the 142 resolved basins with clean footprints. Representation is what changes the answer: handed over as raw 12-month history the same neighbor signal yields nothing – whether as an encoder input channel or as a dedicated correction stage – while as a single state propagated to the

835 forecast target time by the filter it contributes +0.91 to +1.96% at every lead 1–6, beating 20 of 20 degree-matched random-graph placebos in every lead–seed cell, persistent across leads and largest at lead 6. Two controls carry that claim, and neither is the leaderboard: the placebo family, which makes it a statement about the neighbor’s information rather than about the extra stage, and the footprint-sharing stratification, in which the correction is positive in all 60 strata-by-lead cells, is worth +0.89 to +2.00% in the very basins where the linear effect is negative, and is if anything weakest in the small, most footprint-shared
840 basins that the artifact would favour – so the artifact that explains the linear pocket does not explain the correction. What does explain it remains unidentified.

Short leads are a filtering problem; long leads are a forcing problem; the spatial signal is small, invisible to linear methods, and usable only in propagated-state form. The obvious next system combines all three, and the obvious next practice is to benchmark it – and everything else in this literature – against persistence of the filtered state rather than persistence of the raw
845 observation.

Code and data availability. CSR GRACE/GRACE-FO mascon solutions are available from the Center for Space Research (<https://www2.csr.utexas.edu/grace/>). ERA5-Land data are available from the Copernicus Climate Data Store (<https://cds.climate.copernicus.eu/>) (Muñoz-Sabater et al., 2021). Climate indices are available from NOAA PSL (<https://psl.noaa.gov/data/climateindices/>). The GRACE-FCast hindcast is available from PANGAEA (Li, 2025). Analysis code, derived basin tables, and all per-prediction result files sufficient to reproduce every
850 number in this paper will be archived on Zenodo upon acceptance; Appendix A maps the model names used in this paper to the identifiers in those files.

Appendix A: Model names and archived identifiers

The archived result files identify each model by the internal identifier used in the analysis code. Table A1 maps the model names of Table 1 to those identifiers, and Table A2 lists the archived file behind each table and figure; the identifiers appear
855 nowhere else in the paper. Seed suffixes (`_s0`, `_s1`, `_s2`) denote single training seeds and `_ens` the seed ensemble.

Author contributions. I.K. designed the study, wrote the analysis code, performed the experiments and audits, and wrote the manuscript.

Competing interests. The authors declare that they have no conflict of interest.

Acknowledgements. The author thanks the Center for Space Research (University of Texas at Austin) for the open mascon solutions and ancillary files, the Copernicus Climate Change Service for ERA5-Land, NOAA PSL for the climate indices, and Li and Kusche for publishing
860 their gridded hindcast, which made the head-to-head comparison possible. This research received no external funding.

Table A1. Model names used in the paper and the corresponding identifiers in the archived result files.

Paper name	Identifier(s)
Climatology	climatology_zero
Persistence	persistence
Damped persistence	damped_persistence_rho, damped_persistence_reg
Pooled ridge (and climate-index variant)	ridge_own_lags, ridge_own_plus_indices
Per-basin ridge	ridge_own_perbasin
Kalman forecast	kalman_ar1
Ridge on filtered states	kalman_own_ridge
Neighbor ridge (registered contrast)	kalman_corr_top1
Neighbor ridge, distance-restricted	kalman_corr_min300_top1, kalman_corr_min500_top1, kalman_corr_min1000_top1
Neighbor ridge, alternative graphs	kalman_geo_top1, kalman_geo_top2, kalman_geo_top3, kalman_corr_top2, kalman_corr_top3
Own-basin ridge / ERA5 ridge	ridge_own, ridge_own_era5
Neighbor ridge (forcing experiments)	ridge_corr_top1, ridge_corr_top1_era5
GBM head	gbm_own_era5, gbm_corr_top1_era5
MLP head	mlp_own_era5_s0
Residual MLP	resmlp_corr_top1, resmlp_corr_top1_era5, resmlp_own_era5
LSTM (and neighbor-channel variant)	lstm_own, lstm_own_era5, lstm_corr_top1_era5
Stacked system	lstmres_corr_top1
Own-state control	lstmres_own
GAT	gnn_corr_top1, gnn_corr_top2_era5, and variants
GRACE-FCast (full / non-seasonal)	li_lstm_full, li_lstm_nonseas

References

Ahi, G. O. and Cekim, H. O.: Long-term temporal prediction of terrestrial water storage changes over global basins using GRACE and limited GRACE-FO data, *Acta Geodaetica et Geophysica*, 56, 321–344, <https://doi.org/10.1007/s40328-021-00338-4>, 2021.

Ahmed, M., Sultan, M., Elbayoumi, T., and Tissot, P.: Forecasting GRACE Data over the African Watersheds Using Artificial Neural Networks, *Remote Sensing*, 11, 1769, <https://doi.org/10.3390/rs11151769>, 2019.

Arsenault, K. R., Shukla, S., Hazra, A., Getirana, A., McNally, A., Kumar, S. V., Koster, R. D., Peters-Lidard, C. D., Zaitchik, B. F., Badr, H., Jung, H. C., Narapusetty, B., Navari, M., Wang, S., Mocko, D. M., Funk, C., Harrison, L., Husak, G. J., Adoum, A., Galu, G., Magadzire, T., Roningen, J., Shaw, M., Eylander, J., Bergaoui, K., McDonnell, R. A., and Verdin, J. P.: The NASA Hydrological Forecast System for Food and Water Security Applications, *Bulletin of the American Meteorological Society*, 101, E1007–E1025, <https://doi.org/10.1175/BAMS-D-18-0264.1>, 2020.

Table A2. Archived source files for the tables and figures.

Item		Archived source file(s)
Table 2, Fig. 1		results/paper_baseline_ladder.csv; pairwise tests in results/paper_baseline_contrasts.csv
Table 3, Fig. 2		results/phase8b_li_comparison_headline.csv, results/phase8b_li_comparison_perbasin.csv
Table 4, Fig. 3		results/phase3b_summary.csv, results/phase6_era5_headline.csv
Sharing-area (Sect. 4.4)	cross-stratification	results/resolution_cross_2x2.csv, results/resolution_sweep_area.csv, results/resolution_sweep_contamination.csv
Table 5, Fig. 5		results/phase8b_h16_ensemble_headline.csv, results/phase8b_h16_headline.csv, results/phase8b_h16_summary.csv, results/phase8b_lstm_h46_placebo_monthly.csv
Stratification of the stacked correction (Sect. 4.5)		results/phase8_stratification.csv
Forcing and architecture comparisons (Sect. 4.2)		results/phase6_era5_headline.csv, results/phase7_corrected_analysis.md, results/phase6_era5_attribution.csv (and _folds / _continent / _fdr companions), results/flat12_train85_sensitivity.csv
Fig. 4, Fig. 7		results/phase5_perbasin_fdr_h1.csv, results/phase6_basin_analysis_summary.csv, results/phase6_basin_analysis_q4.csv

Arzoumanidis, L., Johannsen, L., Middendorf, K., Eicker, A., and Dehbi, Y.: Reconstructing GRACE Terrestrial Water Storage with Spatio-Temporal Graph Neural Networks: An Application to South America, <https://arxiv.org/abs/2606.23833>, 2026.

Benjamini, Y. and Hochberg, Y.: Controlling the False Discovery Rate: A Practical and Powerful Approach to Multiple Testing, *Journal of the Royal Statistical Society Series B: Statistical Methodology*, 57, 289–300, <https://doi.org/10.1111/j.2517-6161.1995.tb02031.x>, 1995.

875 Diebold, F. X. and Mariano, R. S.: Comparing Predictive Accuracy, *Journal of Business & Economic Statistics*, 13, 253–263, <https://doi.org/10.1080/07350015.1995.10524599>, 1995.

Getirana, A., Rodell, M., Kumar, S., Beaudoin, H. K., Arsenault, K., Zaitchik, B., Save, H., and Bettadpur, S.: GRACE Improves Seasonal Groundwater Forecast Initialization over the United States, *Journal of Hydrometeorology*, 21, 59–71, <https://doi.org/10.1175/JHM-D-19-0096.1>, 2020.

- 880 Harvey, D., Leybourne, S., and Newbold, P.: Testing the equality of prediction mean squared errors, *International Journal of Forecasting*, 13, 281–291, [https://doi.org/10.1016/S0169-2070\(96\)00719-4](https://doi.org/10.1016/S0169-2070(96)00719-4), 1997.
- Hochreiter, S. and Schmidhuber, J.: Long Short-Term Memory, *Neural Computation*, 9, 1735–1780, <https://doi.org/10.1162/neco.1997.9.8.1735>, 1997.
- Humphrey, V. and Gudmundsson, L.: GRACE-REC: a reconstruction of climate-driven water storage changes over the last century, *Earth System Science Data*, 11, 1153–1170, <https://doi.org/10.5194/essd-11-1153-2019>, 2019.
- 885 Kalman, R. E.: A New Approach to Linear Filtering and Prediction Problems, *Journal of Basic Engineering*, 82, 35–45, <https://doi.org/10.1115/1.3662552>, 1960.
- Kankanige, D., Vishwakarma, B. D., Liu, Y. Y., and Sharma, A.: Can a Persistence-Based Prediction of GRACE Gravity Field Coefficients Reduce Uncertainties at Finer Temporal Scales?, *Earth and Space Science*, 13, <https://doi.org/10.1029/2025EA004554>, 2026.
- 890 Koster, R. D., Mahanama, S. P. P., Yamada, T. J., Balsamo, G., Berg, A. A., Boisserie, M., Dirmeyer, P. A., Doblas-Reyes, F. J., Drewitt, G., Gordon, C. T., Guo, Z., Jeong, J.-H., Lawrence, D. M., Lee, W.-S., Li, Z., Luo, L., Malyshev, S., Merryfield, W. J., Seneviratne, S. I., Stanelle, T., van den Hurk, B. J. J. M., Vitart, F., and Wood, E. F.: Contribution of land surface initialization to subseasonal forecast skill: First results from a multi-model experiment, *Geophysical Research Letters*, 37, <https://doi.org/10.1029/2009GL041677>, 2010.
- Künsch, H. R.: The Jackknife and the Bootstrap for General Stationary Observations, *The Annals of Statistics*, 17, 1217–1241, <https://doi.org/10.1214/aos/1176347265>, 1989.
- 895 Landerer, F. W., Flechtner, F. M., Save, H., Webb, F. H., Bandikova, T., Bertiger, W. I., Bettadpur, S. V., Byun, S. H., Dahle, C., Dobsław, H., Fahnestock, E., Harvey, N., Kang, Z., Kruizinga, G. L. H., Loomis, B. D., McCullough, C., Murböck, M., Nagel, P., Paik, M., Pie, N., Poole, S., Strelakov, D., Tamisiea, M. E., Wang, F., Watkins, M. M., Wen, H.-Y., Wiese, D. N., and Yuan, D.-N.: Extending the Global Mass Change Data Record: GRACE Follow-On Instrument and Science Data Performance, *Geophysical Research Letters*, 47, <https://doi.org/10.1029/2020GL088306>, 2020.
- 900 Lehner, B. and Grill, G.: Global river hydrography and network routing: baseline data and new approaches to study the world’s large river systems, *Hydrological Processes*, 27, 2171–2186, <https://doi.org/10.1002/hyp.9740>, 2013.
- Li, B., Hazra, A., McNally, A., Slinski, K., Shukla, S., and Anderson, W.: Skills in sub-seasonal to seasonal terrestrial water storage forecasting: insights from the FEWS NET land data assimilation system, *Hydrology and Earth System Sciences*, 30, 1097–1115, <https://doi.org/10.5194/hess-30-1097-2026>, 2026.
- 905 Li, F.: A Global Dataset for Seasonal to Annual Forecasts of GRACE-like Gridded Terrestrial Water Storage (2010–2024), <https://doi.org/10.1594/PANGAEA.973113>, 2025.
- Li, F. and Kusche, J.: Observation-Driven Forecast of Global Terrestrial Water Storage and Evaluation for 2010–2024, *Water Resources Research*, 62, <https://doi.org/10.1029/2025WR041710>, 2026.
- 910 Li, F., Kusche, J., Rietbroek, R., Wang, Z., Forootan, E., Schulze, K., and Lück, C.: Comparison of Data-Driven Techniques to Reconstruct (1992–2002) and Predict (2017–2018) GRACE-Like Gridded Total Water Storage Changes Using Climate Inputs, *Water Resources Research*, 56, <https://doi.org/10.1029/2019WR026551>, 2020.
- Li, F., Kusche, J., Sneeuw, N., Siebert, S., Gerdener, H., Wang, Z., Chao, N., Chen, G., and Tian, K.: Forecasting Next Year’s Global Land Water Storage Using GRACE Data, *Geophysical Research Letters*, 51, <https://doi.org/10.1029/2024GL109101>, 2024.
- 915 Li, F., Springer, A., Kusche, J., Gutknecht, B. D., and Ewerdwalbesloh, Y.: Reanalysis and Forecasting of Total Water Storage and Hydrological States by Combining Machine Learning With CLM Model Simulations and GRACE Data Assimilation, *Water Resources Research*, 61, <https://doi.org/10.1029/2024WR037926>, 2025.

Mo, S., Zhong, Y., Forootan, E., Mehrnegar, N., Yin, X., Wu, J., Feng, W., and Shi, X.: Bayesian convolutional neural networks for predicting the terrestrial water storage anomalies during GRACE and GRACE-FO gap, *Journal of Hydrology*, 604, 127 244, <https://doi.org/10.1016/j.jhydrol.2021.127244>, 2022.

Mo, S., Schumacher, M., van Dijk, A. I. J. M., Shi, X., Wu, J., and Forootan, E.: Near-Real-Time Monitoring of Global Terrestrial Water Storage Anomalies and Hydrological Droughts, *Geophysical Research Letters*, 52, <https://doi.org/10.1029/2024GL112677>, 2025.

Muñoz-Sabater, J., Dutra, E., Agustí-Panareda, A., Albergel, C., Arduini, G., Balsamo, G., Boussetta, S., Choulga, M., Harrigan, S., Hersbach, H., Martens, B., Miralles, D. G., Piles, M., Rodríguez-Fernández, N. J., Zsoter, E., Buontempo, C., and Thépaut, J.-N.: ERA5-Land: a state-of-the-art global reanalysis dataset for land applications, *Earth System Science Data*, 13, 4349–4383, <https://doi.org/10.5194/essd-13-4349-2021>, 2021.

Newey, W. K. and West, K. D.: A Simple, Positive Semi-Definite, Heteroskedasticity and Autocorrelation Consistent Covariance Matrix, *Econometrica*, 55, 703–708, <https://doi.org/10.2307/1913610>, 1987.

Nie, W., Kumar, S. V., Chen, J., Zhao, L., Skulovich, O., Yoo, J., Pflug, J., Ahmad, S. K., and Konapala, G.: Rethinking deep learning: linear regression remains a key benchmark in predicting terrestrial water storage, <https://arxiv.org/abs/2510.10799>, 2025.

Niraula, B. and Goessling, H. F.: Spatial Damped Anomaly Persistence of the Sea Ice Edge as a Benchmark for Dynamical Forecast Systems, *Journal of Geophysical Research: Oceans*, 126, <https://doi.org/10.1029/2021JC017784>, 2021.

Pechlivanidis, I. G., Crochemore, L., Rosberg, J., and Bosshard, T.: What Are the Key Drivers Controlling the Quality of Seasonal Streamflow Forecasts?, *Water Resources Research*, 56, <https://doi.org/10.1029/2019WR026987>, 2020.

Pezeshki, M., Kaba, S.-O., Bengio, Y., Courville, A., Precup, D., and Lajoie, G.: Gradient Starvation: A Learning Proclivity in Neural Networks, <https://arxiv.org/abs/2011.09468>, *advances in Neural Information Processing Systems* 34, 2021.

Reager, J. T., Thomas, B. F., and Famiglietti, J. S.: River basin flood potential inferred using GRACE gravity observations at several months lead time, *Nature Geoscience*, 7, 588–592, <https://doi.org/10.1038/ngeo2203>, 2014.

Save, H., Bettadpur, S., and Tapley, B. D.: High-resolution CSR GRACE RL05 mascons, *Journal of Geophysical Research: Solid Earth*, 121, 7547–7569, <https://doi.org/10.1002/2016JB013007>, 2016.

Schreiber, T. and Schmitz, A.: Improved Surrogate Data for Nonlinearity Tests, *Physical Review Letters*, 77, 635–638, <https://doi.org/10.1103/PhysRevLett.77.635>, 1996.

Schreiber, T. and Schmitz, A.: Surrogate time series, *Physica D: Nonlinear Phenomena*, 142, 346–382, [https://doi.org/10.1016/S0167-2789\(00\)00043-9](https://doi.org/10.1016/S0167-2789(00)00043-9), 2000.

Shukla, S. and Lettenmaier, D. P.: Seasonal hydrologic prediction in the United States: understanding the role of initial hydrologic conditions and seasonal climate forecast skill, *Hydrology and Earth System Sciences*, 15, 3529–3538, <https://doi.org/10.5194/hess-15-3529-2011>, 2011.

Springer, A., De Lannoy, G., Rodell, M., Ewerdtwalbesloh, Y., Gerdener, H., Khaki, M., Li, B., Li, F., Schumacher, M., Tangdamrongsub, N., Tourian, M. J., Nie, W., and Kusche, J.: A review of current best practices and future directions in assimilating GRACE/FO terrestrial water storage data into numerical models, *Hydrology and Earth System Sciences*, 30, 985–1022, <https://doi.org/10.5194/hess-30-985-2026>, 2026.

Steidl, V. and Zhu, X. X.: Hierarchical Graph Networks for Forecasting Terrestrial Water Storage Anomalies, in: *Machine Learning and the Physical Sciences Workshop, 39th Conference on Neural Information Processing Systems (NeurIPS)*, <https://ml4physicalsciences.github.io/2025/>, 2025.

- 955 Tapley, B. D., Bettadpur, S., Watkins, M., and Reigber, C.: The gravity recovery and climate experiment: Mission overview and early results, *Geophysical Research Letters*, 31, <https://doi.org/10.1029/2004GL019920>, 2004.
- Veličković, P., Cucurull, G., Casanova, A., Romero, A., Liò, P., and Bengio, Y.: Graph Attention Networks, <https://arxiv.org/abs/1710.10903>, international Conference on Learning Representations, 2018.
- Wilks, D. S.: “The Stippling Shows Statistically Significant Grid Points”: How Research Results are Routinely Overstated and Overinter-
 960 preted, and What to Do about It, *Bulletin of the American Meteorological Society*, 97, 2263–2273, <https://doi.org/10.1175/BAMS-D-15-00267.1>, 2016.
- Wood, A. W. and Lettenmaier, D. P.: An ensemble approach for attribution of hydrologic prediction uncertainty, *Geophysical Research Letters*, 35, <https://doi.org/10.1029/2008GL034648>, 2008.
- Yin, J., Slater, L. J., Khouakhi, A., Yu, L., Liu, P., Li, F., Pokhrel, Y., and Gentile, P.: GTWS-MLrec: global terrestrial water storage
 965 reconstruction by machine learning from 1940 to present, *Earth System Science Data*, 15, 5597–5615, <https://doi.org/10.5194/essd-15-5597-2023>, 2023.
- Yossef, N. C., Winsemius, H., Weerts, A., van Beek, R., and Bierkens, M. F. P.: Skill of a global seasonal streamflow forecasting system, relative roles of initial conditions and meteorological forcing, *Water Resources Research*, 49, 4687–4699, <https://doi.org/10.1002/wrcr.20350>, 2013.
- 970 Zhang, L., Shen, Y., Ji, K., Wang, F., and Chen, Q.: An integrated postprocessing approach for extracting time variable signals from GRACE monthly gravity field models, *Journal of Hydrology*, 661, 133 552, <https://doi.org/10.1016/j.jhydrol.2025.133552>, 2025.
- Zhu, E., Yuan, X., and Wood, A. W.: Benchmark decadal forecast skill for terrestrial water storage estimated by an elasticity framework, *Nature Communications*, 10, 1237, <https://doi.org/10.1038/s41467-019-09245-3>, 2019.